

# Time Series Forecasting report for total industry

Kevork Sulahian

2021-04-22

```
library(readxl)
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```

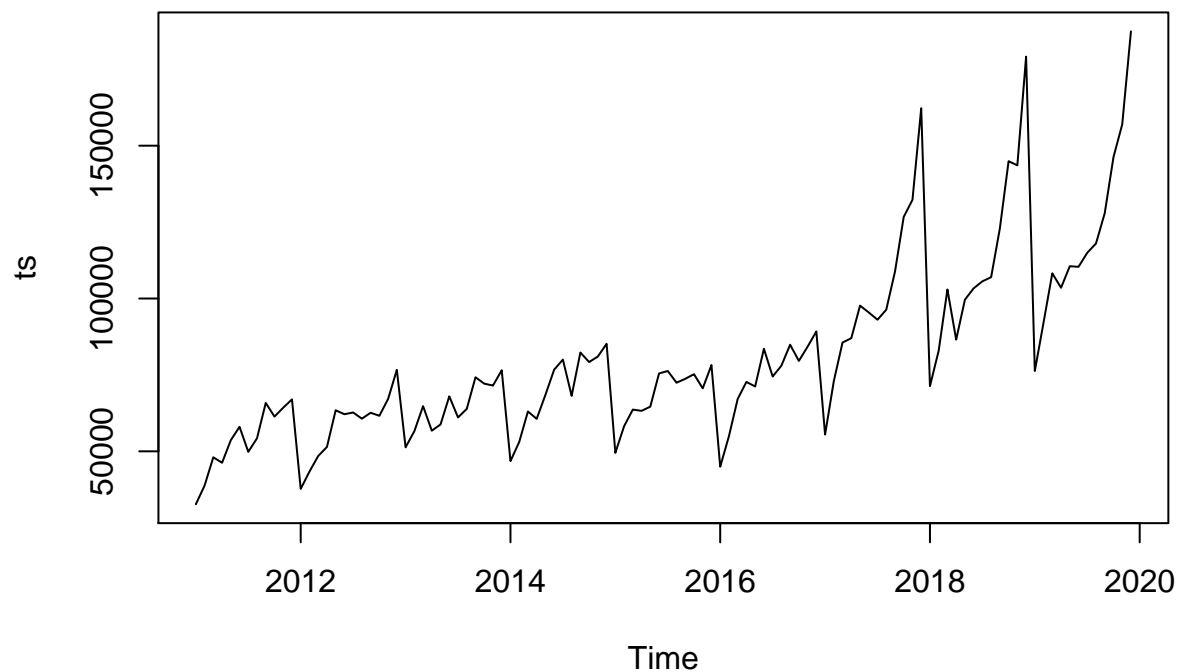
```
df <- read_xls('economy.xls', sheet='2011-2019 NACE 2')
```

```
## New names:
## * ' ' -> ...2
## * ' ' -> ...3
## * ' ' -> ...4
## * ' ' -> ...5
## * ' ' -> ...6
## * ...
```

```
df = df[11,]
df = df[-c(1,3)]
rownames(df) = df[1]
```

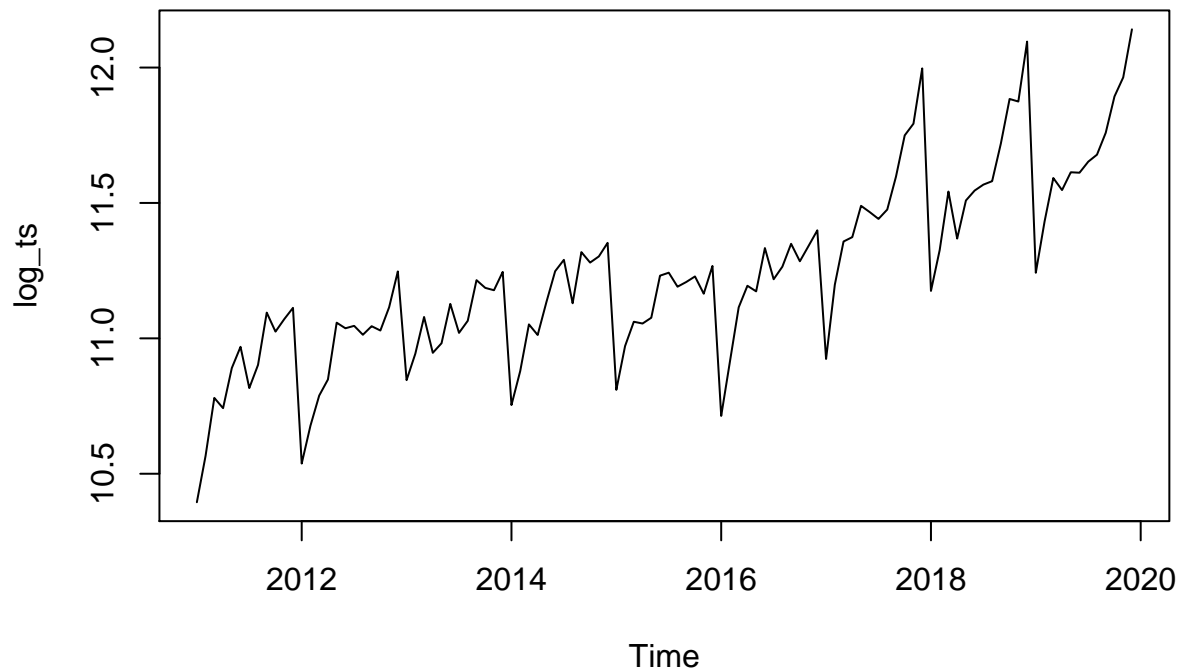
```
## Warning: Setting row names on a tibble is deprecated.
```

```
df = df[-1]
df = t(df)
df[110] = "155770.7"
df[] <- sapply(df[],function(x) as.numeric(as.character(x)))
df = as.numeric(df)
# df= df * 1000000
df = df[-c(109:120)]
ts = ts(df, start = c(2011,1), frequency = c(12))
```



In this case, it appears that an additive model is not appropriate for describing this time series, since the size of the seasonal fluctuations and random fluctuations seem to increase with the level of the time series. Thus, we may need to transform the time series in order to get a transformed time series that can be described using an additive model. For example, we can transform the time series by calculating the natural log of the original data:

```
log_ts <- log(ts)
plot.ts(log_ts)
```



## ##Decomposing Time Series

Decomposing a time series means separating it into its constituent components, which are usually a trend component and an irregular component, and if it is a seasonal time series, a seasonal component.

###Decomposing Seasonal Data A seasonal time series consists of a trend component, a seasonal component and an irregular component. Decomposing the time series means separating the time series into these three components: that is, estimating these three components.

```
ts_components <- decompose(ts)
```

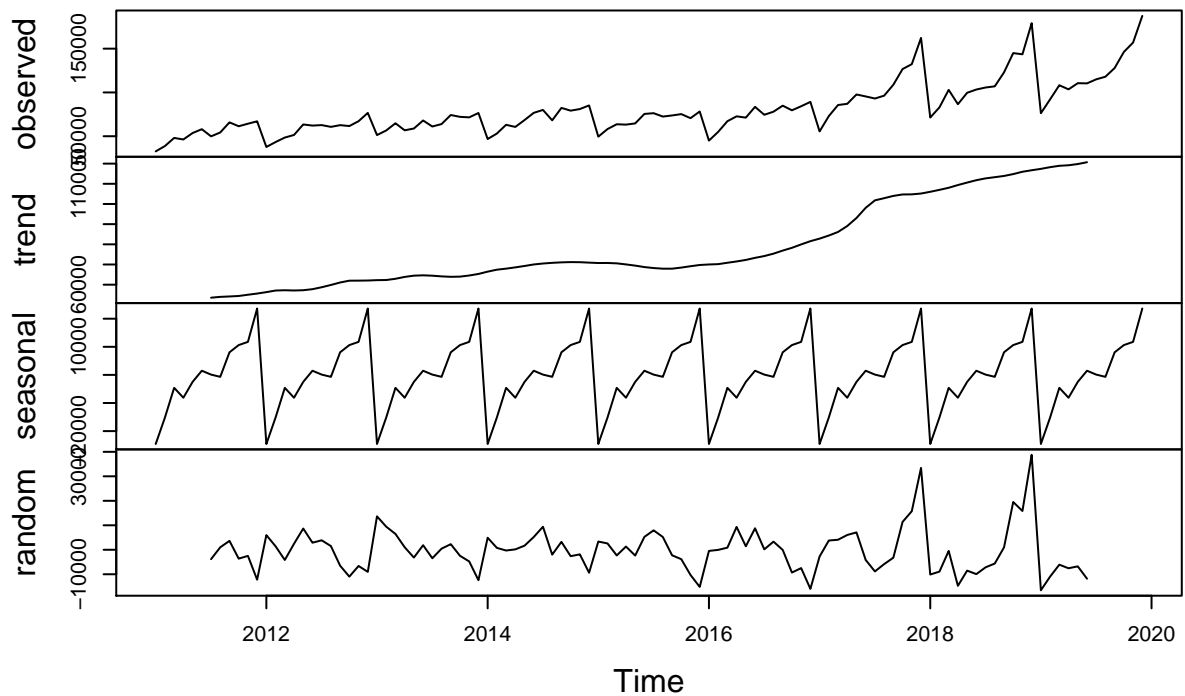
we can print out the estimated values of the seasonal component

```
ts_components$seasonal
```

##		Jan	Feb	Mar	Apr	May	Jun
##	2011	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2012	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2013	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2014	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2015	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2016	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2017	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2018	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##	2019	-24588.4065	-15153.7133	-4601.1352	-8123.8237	-2480.5076	1488.2299
##		Jul	Aug	Sep	Oct	Nov	Dec

## 2011	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2012	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2013	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2014	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2015	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2016	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2017	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2018	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622
## 2019	109.2596	-686.7487	8045.2492	10616.0039	11749.1299	23626.4622

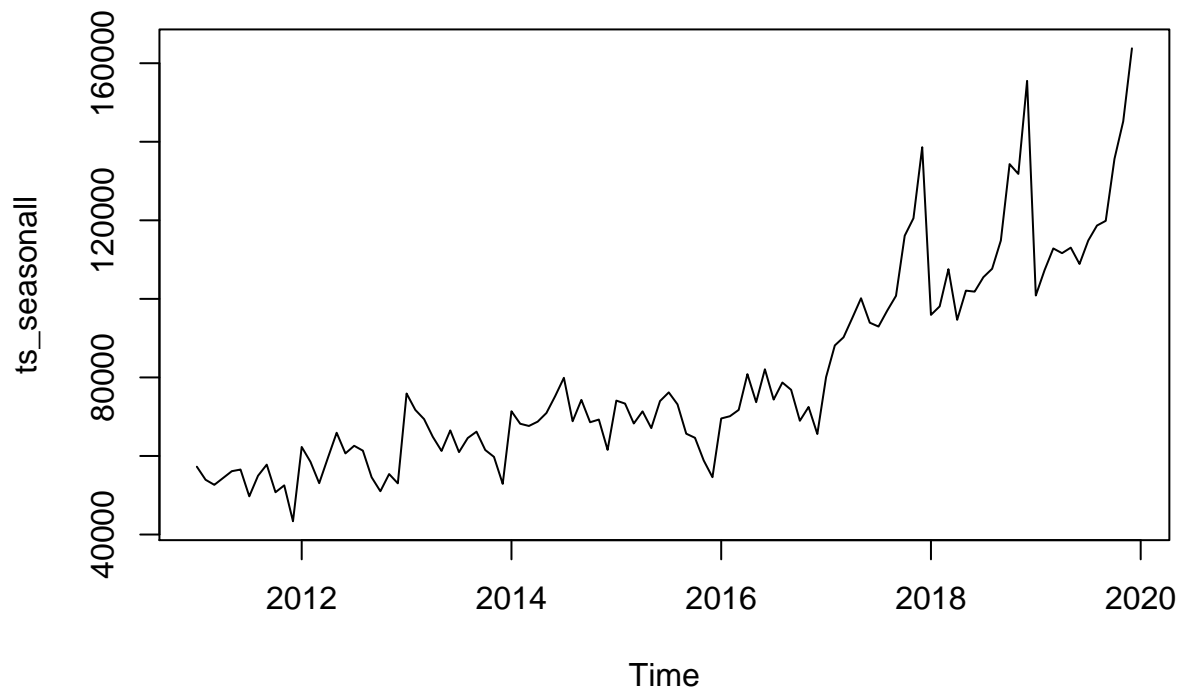
## Decomposition of additive time series



The plot above shows the original time series (top), the estimated trend component (second from top), the estimated seasonal component (third from top), and the estimated irregular component (bottom)

## Seasonally Adjusting

```
ts_seasonall <- ts - ts_components$seasonal
```



```
## Holt-Winters Exponential Smoothing
```

```
ts_forcaste <- HoltWinters(ts)
ts_forcaste
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = ts)
##
## Smoothing parameters:
##  alpha: 0.2187957
##  beta : 0
##  gamma: 1
##
## Coefficients:
##           [,1]
## a    116063.0702
## b      412.1641
## s1  -38745.2509
## s2  -23525.0504
## s3   -6821.8331
## s4  -12611.9619
## s5   -4592.2182
## s6   -3071.9176
## s7    2147.6401
```

```
## s8      5161.0306
## s9     16027.7968
## s10    35096.6029
## s11    43103.5368
## s12    71346.4298
```

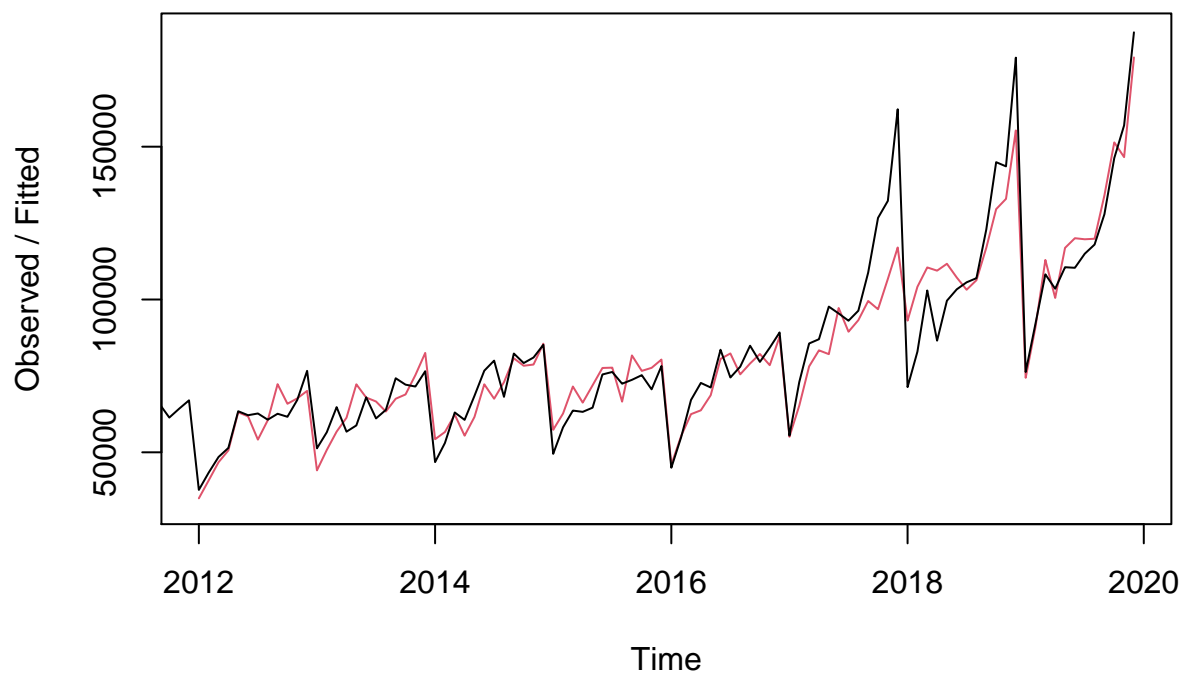
```
#
```

The value of alpha (0.41) is relatively low, indicating that the estimate of the level at the current time point is based upon both recent observations and some observations in the more distant past. The value of beta is 0.00, indicating that the estimate of the slope  $b$  of the trend component is not updated over the time series, and instead is set equal to its initial value. This makes good intuitive sense, as the level changes quite a bit over the time series, but the slope  $b$  of the trend component remains roughly the same. In contrast, the value of gamma (0.96) is high, indicating that the estimate of the seasonal component at the current time point is just based upon very recent observations

```
ts_forcaste$SSE
```

```
## [1] 8847392124
```

## Holt-Winters filtering

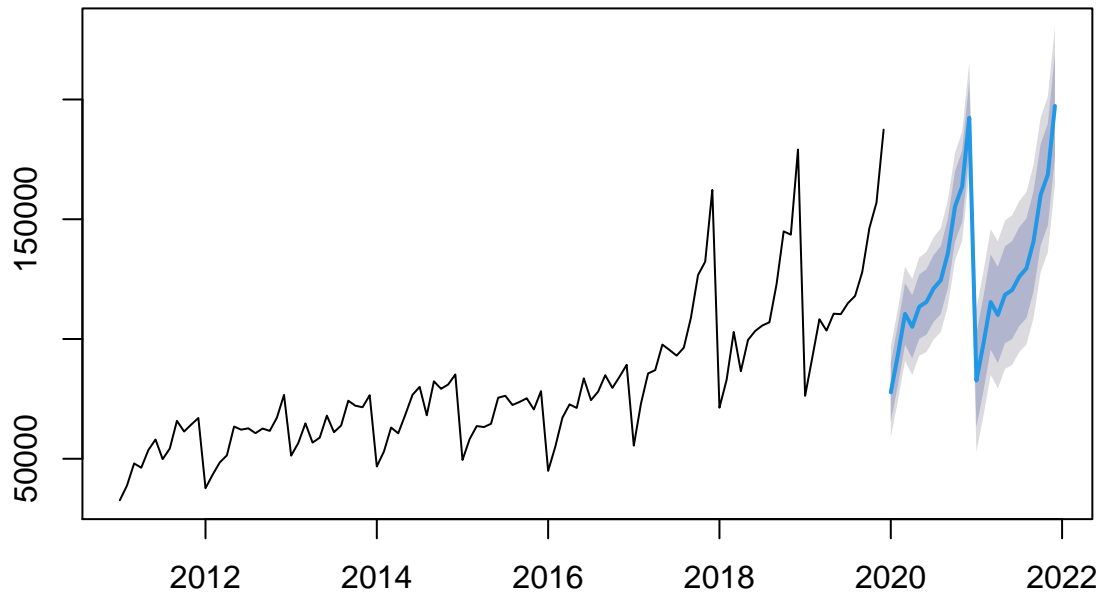


```
ts_forcaste2 = forecast::forecast.HoltWinters(ts_forcaste, h= 24)
(as.data.frame(ts_forcaste2))[1]
```

```
##          Point Forecast
```

## Jan 2020	77729.98
## Feb 2020	93362.35
## Mar 2020	110477.73
## Apr 2020	105099.76
## May 2020	113531.67
## Jun 2020	115464.14
## Jul 2020	121095.86
## Aug 2020	124521.41
## Sep 2020	135800.34
## Oct 2020	155281.31
## Nov 2020	163700.41
## Dec 2020	192355.47
## Jan 2021	82675.95
## Feb 2021	98308.32
## Mar 2021	115423.70
## Apr 2021	110045.73
## May 2021	118477.64
## Jun 2021	120410.11
## Jul 2021	126041.83
## Aug 2021	129467.38
## Sep 2021	140746.31
## Oct 2021	160227.28
## Nov 2021	168646.38
## Dec 2021	197301.44

### Forecasts from HoltWinters



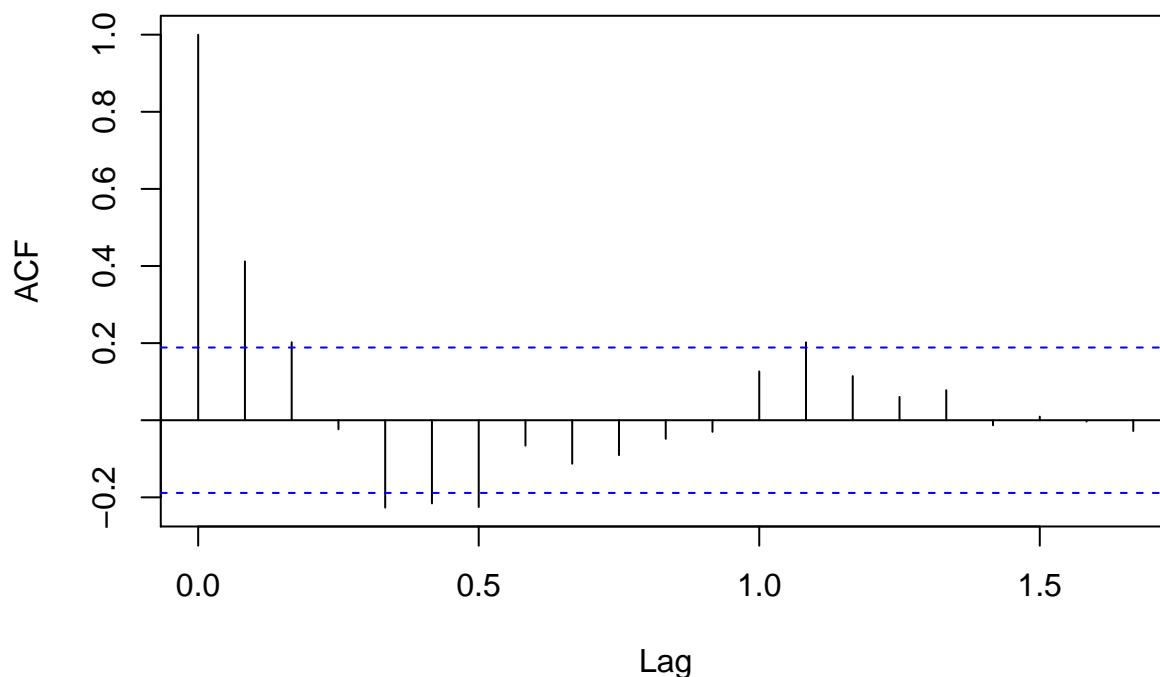
## Growth

```
year_2019 <- window(ts, 2019)
year_2020 <- (as.data.frame(ts_forcaste2))[1][c(1:12),]
year_2021 <- (as.data.frame(ts_forcaste2))[1][c(13:24),]

growth_HW_21 <- growth(sum(year_2021),sum(year_2020))
growth_HW_20 <- growth(sum(year_2020),sum(year_2019))
```

We can investigate whether the predictive model can be improved upon by checking whether the in-sample forecast errors show non-zero autocorrelations at lags 1-20, by making a correlogram and carrying out the

### Series `ts_forcaste2$residuals`



Ljung-Box test:

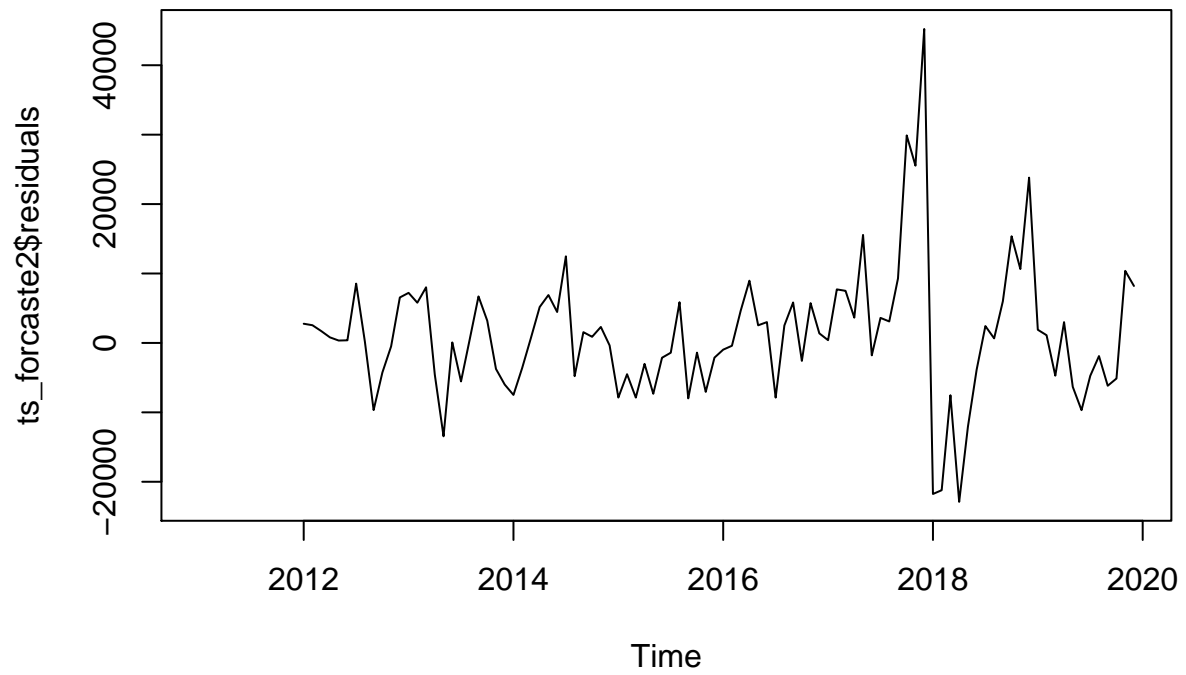
```
##
##  Box-Ljung test
##
## data:  ts_forcaste2$residuals
## X-squared = 48.555, df = 20, p-value = 0.0003554
```

The correlogram shows that the autocorrelations for the in-sample forecast errors do not exceed the significance bounds for lags 1-20. Furthermore, the p-value for Ljung-Box test is 0.2, indicating that there is little evidence of non-zero autocorrelations at lags 1-20.

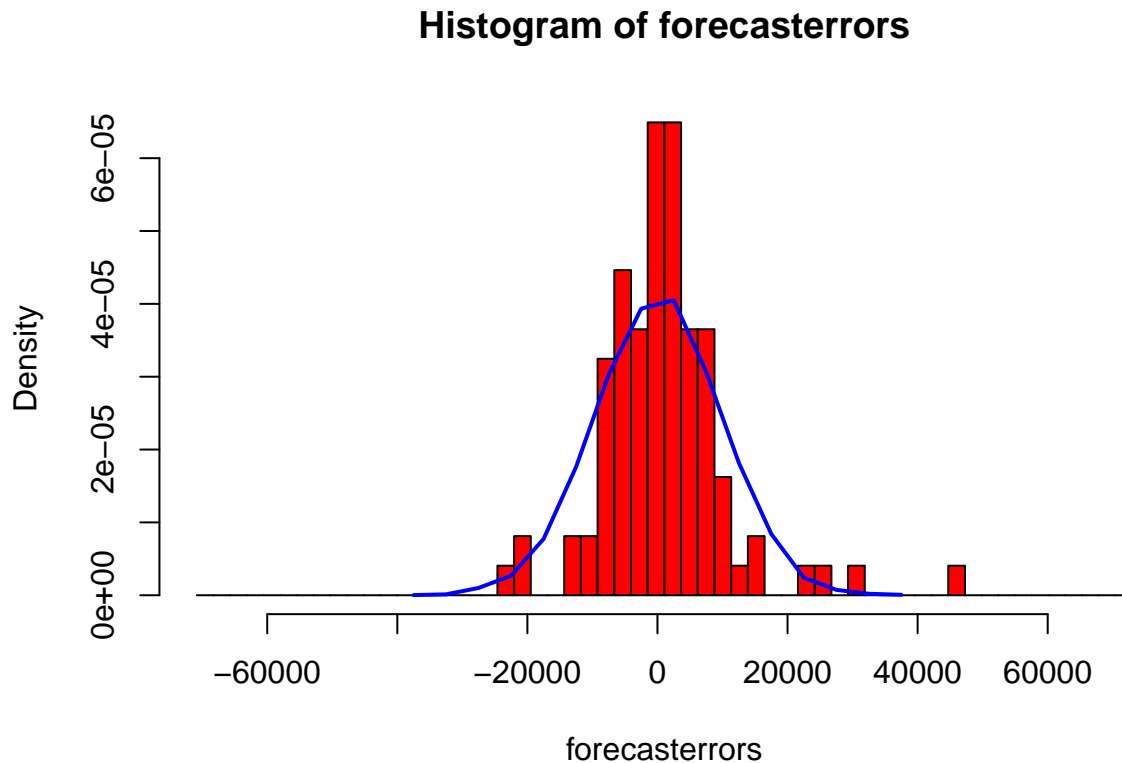
We can check whether the forecast errors have constant variance over time, and are normally distributed with mean zero, by making a time plot of the forecast errors and a histogram (with overlaid normal curve):



```
plot.ts(ts_forcaste2$residuals)
```



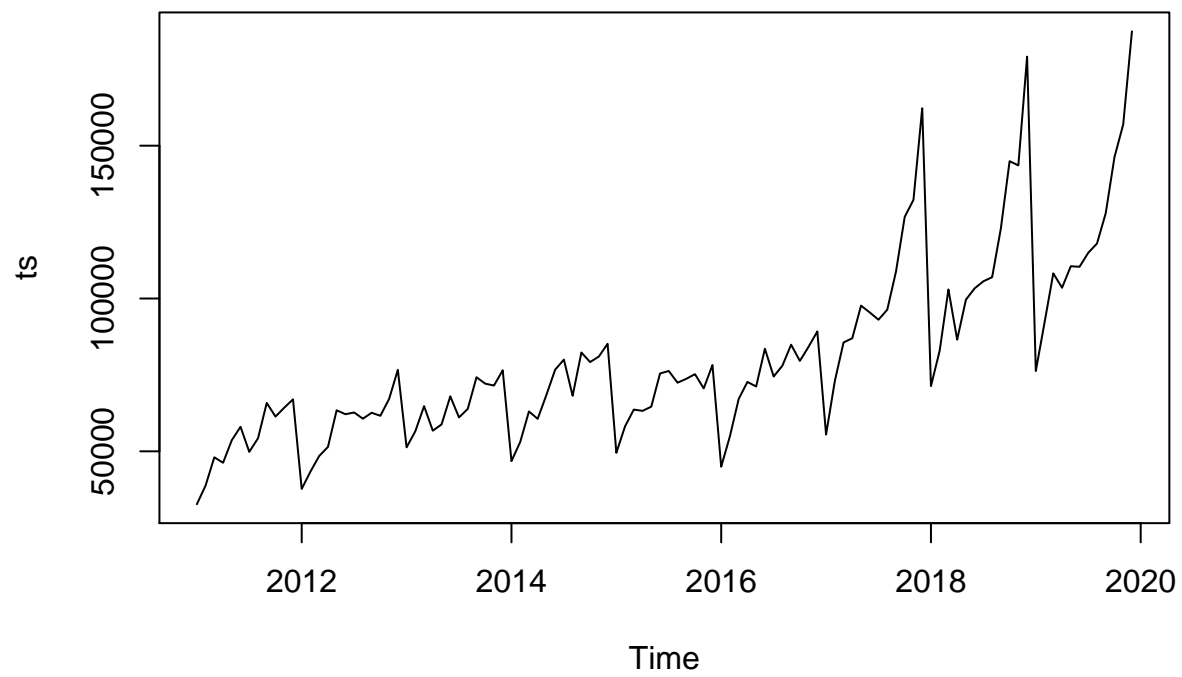
```
plotForecastErrors(ts_forcaste2$residuals)
```



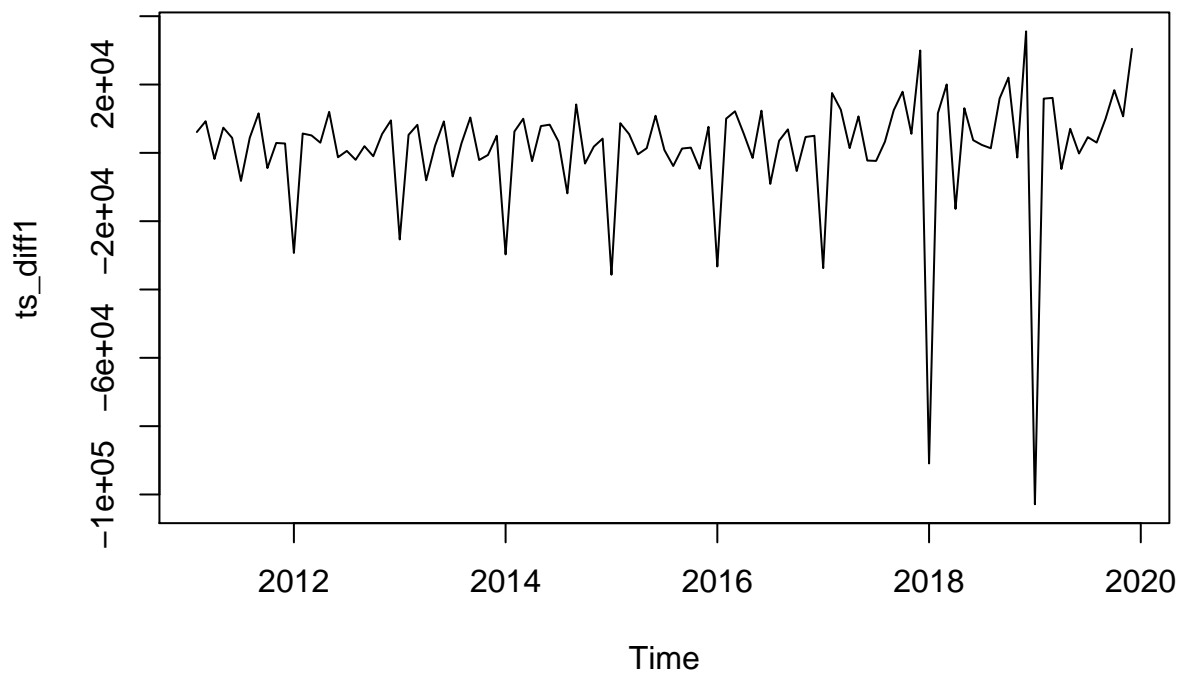
From the time plot, it appears plausible that the forecast errors have constant variance over time. From the histogram of forecast errors, it seems plausible that the forecast errors are normally distributed with mean zero.

Thus, there is little evidence of autocorrelation at lags 1-20 for the forecast errors, and the forecast errors appear to be normally distributed with mean zero and constant variance over time. This suggests that Holt-Winters exponential smoothing provides an adequate predictive model of the log of total productivity, which probably cannot be improved upon. Furthermore, the assumptions upon which the prediction intervals were based are probably valid.

```
plot.ts(ts)
```

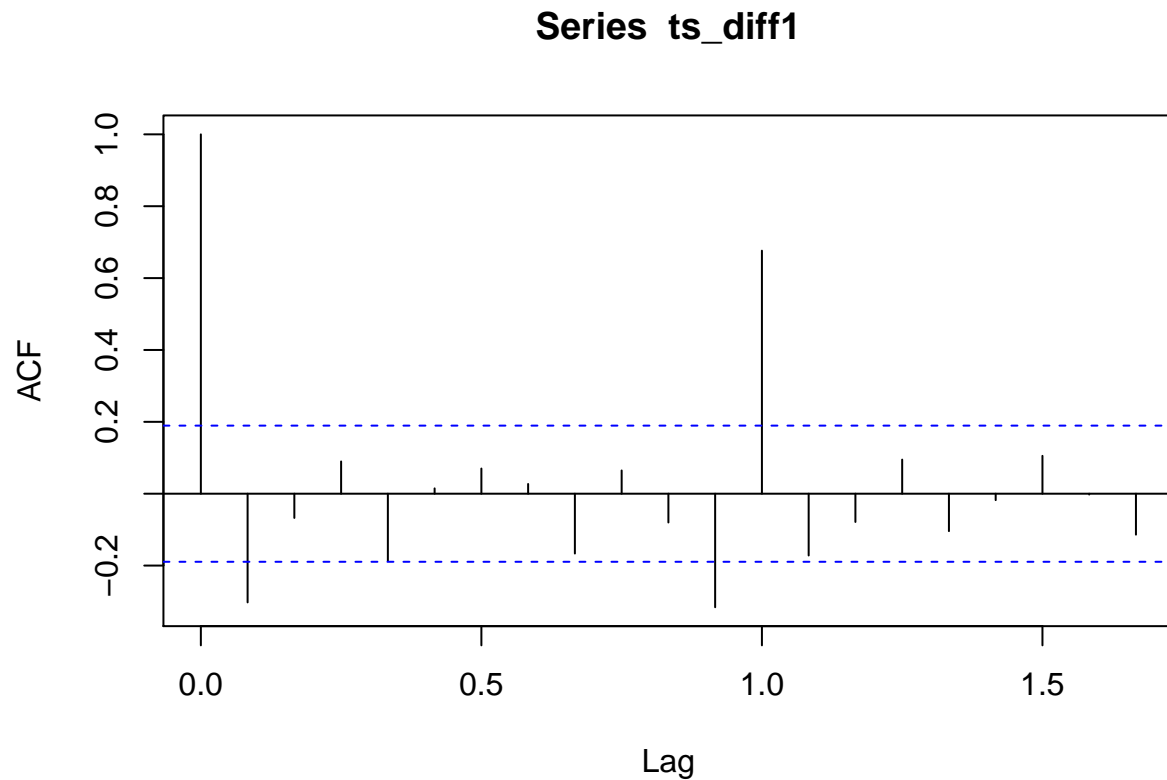


```
ts_diff1 <- diff(ts, differences = 1)
plot.ts(ts_diff1)
```



The time series of differences (above) does appear to be stationary in mean and variance, as the level of the series stays roughly constant over time, and the variance of the series appears roughly constant over time

```
acf(ts_diff1, lag.max=20) # plot a correlogram
```



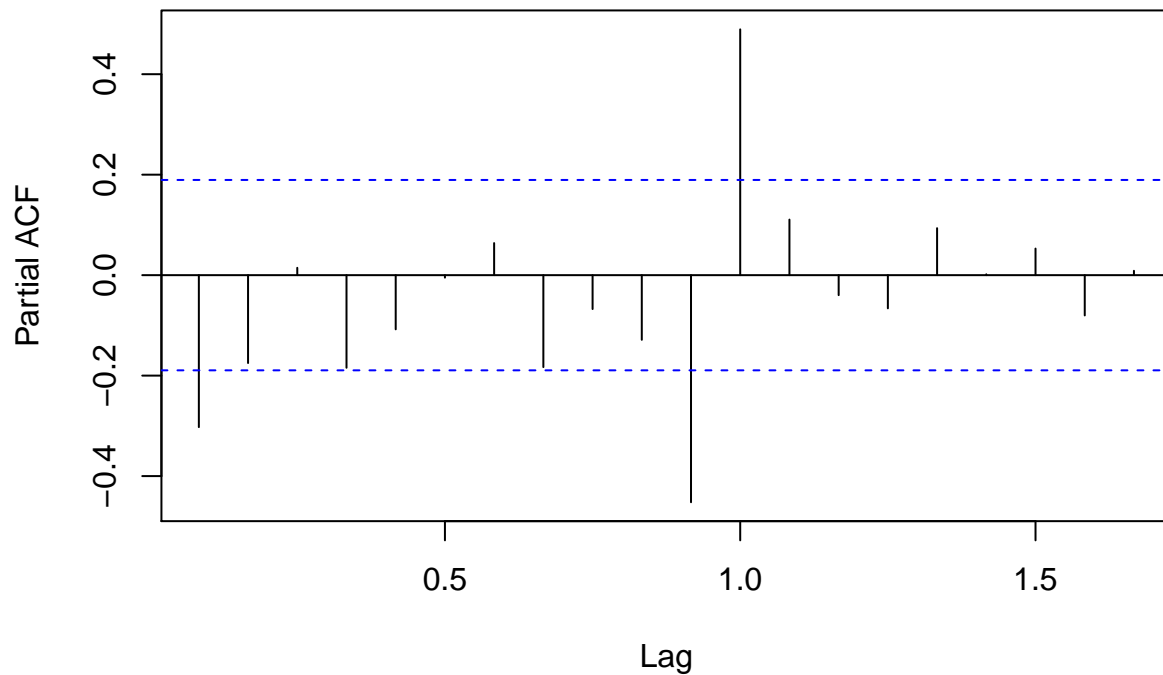
We see from the correlogram that the autocorrelation exceeds the significance bound 3 times but all the others do not exceed

```
acf(ts_diff1, lag.max=20, plot=FALSE) # get the autocorrelation values
```

```
##
## Autocorrelations of series 'ts_diff1', by lag
##
## 0.0000 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333
## 1.000 -0.303 -0.067 0.090 -0.187 0.015 0.070 0.027 -0.167 0.065 -0.080
## 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
## -0.316 0.676 -0.172 -0.079 0.095 -0.104 -0.018 0.105 -0.002 -0.114
```

```
pacf(ts_diff1, lag.max=20) # plot a partial correlogram
```

## Series ts\_diff1



```
pacf(ts_diff1, lag.max=20, plot=FALSE) # get the partial autocorrelation values
```

```
##
## Partial autocorrelations of series 'ts_diff1', by lag
##
## 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333 0.9167
## -0.303 -0.175  0.015 -0.184 -0.108 -0.005  0.064 -0.183 -0.068 -0.129 -0.452
## 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
##  0.489  0.111 -0.040 -0.066  0.094  0.002  0.053 -0.080  0.009
```

## Arima, 1,1,1

```
ts_arima = Arima(ts, order=c(1,1,1),seasonal = list(order = c(1,1,1)))
ts_arima
```

```
## Series: ts
## ARIMA(1,1,1)(1,1,1)[12]
##
## Coefficients:
##          ar1          ma1          sar1          sma1
##          0.6177 -0.9530 -0.8754  0.9998
## s.e.  0.1160  0.0637  0.1033  0.1453
```

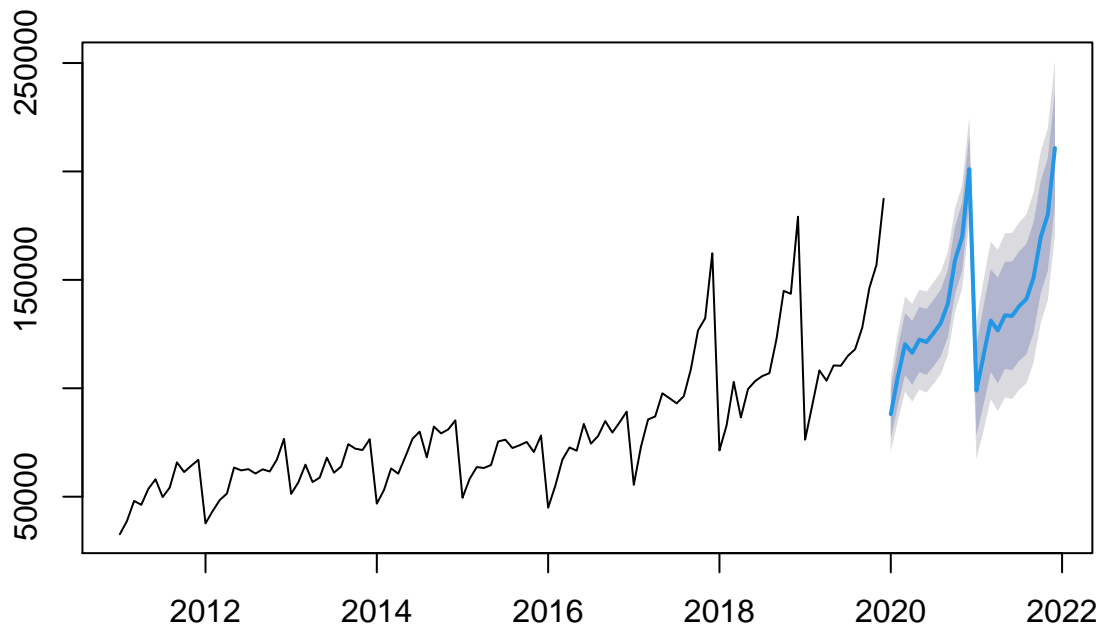
```
##
## sigma^2 estimated as 72394542: log likelihood=-995.44
## AIC=2000.87 AICc=2001.55 BIC=2013.64
```

```
ts_arma_forecast = forecast(ts_arma,h = 24)
ts_arma_forecast
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2020	88039.87	76896.84	99182.9	70998.08	105081.7
## Feb 2020	105436.20	92059.69	118812.7	84978.60	125893.8
## Mar 2020	120324.97	106011.07	134638.9	98433.75	142216.2
## Apr 2020	116361.94	101584.89	131139.0	93762.39	138961.5
## May 2020	122438.66	107400.40	137476.9	99439.62	145437.7
## Jun 2020	121352.80	106147.59	136558.0	98098.44	144607.2
## Jul 2020	125469.91	110145.13	140794.7	102032.68	148907.2
## Aug 2020	130059.01	114639.96	145478.1	106477.61	153640.4
## Sep 2020	139121.13	123622.13	154620.1	115417.45	162824.8
## Oct 2020	159073.19	143502.80	174643.6	135260.34	182886.0
## Nov 2020	169793.04	154156.94	185429.1	145879.69	193706.4
## Dec 2020	201056.74	185359.63	216753.8	177050.08	225063.4
## Jan 2021	98936.41	78260.52	119612.3	67315.37	130557.4
## Feb 2021	115187.92	92475.13	137900.7	80451.71	149924.1
## Mar 2021	131238.08	107501.18	154975.0	94935.63	167540.5
## Apr 2021	126680.82	102348.39	151013.2	89467.59	163894.0
## May 2021	133651.21	108925.06	158377.4	95835.83	171466.6
## Jun 2021	133379.82	108362.67	158397.0	95119.38	171640.3
## Jul 2021	137932.36	112679.79	163184.9	99311.88	176552.8
## Aug 2021	141150.30	115693.91	166606.7	102218.11	180082.5
## Sep 2021	150977.24	125335.73	176618.8	111761.93	190192.5
## Oct 2021	169539.36	143724.16	195354.6	130058.42	209020.3
## Nov 2021	180236.21	154254.57	206217.8	140500.72	219971.7
## Dec 2021	210789.18	184645.87	236932.5	170806.44	250771.9

```
forecast::plot.forecast(ts_arma_forecast)
```

## Forecasts from ARIMA(1,1,1)(1,1,1)[12]



## Growth

```
year_2020 <- (as.data.frame(ts_arma_forecast))[1][c(1:12),]  
year_2021 <- (as.data.frame(ts_arma_forecast))[1][c(13:24),]  
  
growth_ARIMA_21 <- growth(sum(year_2021), sum(year_2020))  
growth_ARIMA_20 <- growth(sum(year_2020), sum(year_2019))
```

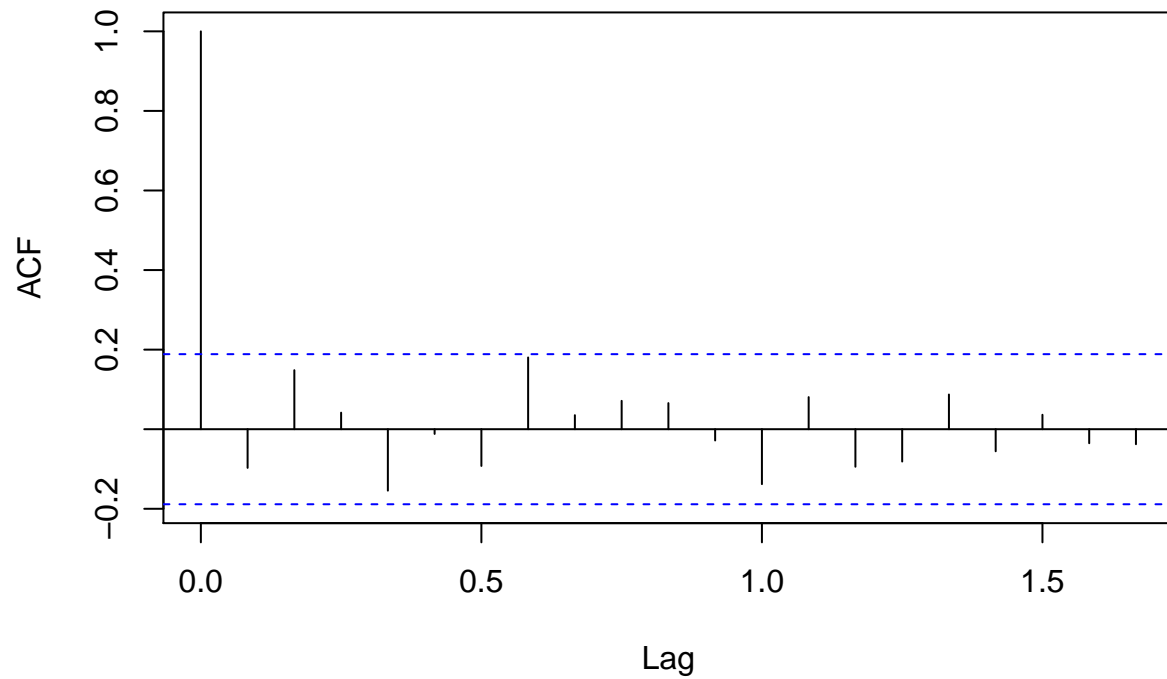
As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arma_forecast$residuals, lag.max=20)
```



### Series ts\_arma\_forecast\$residuals

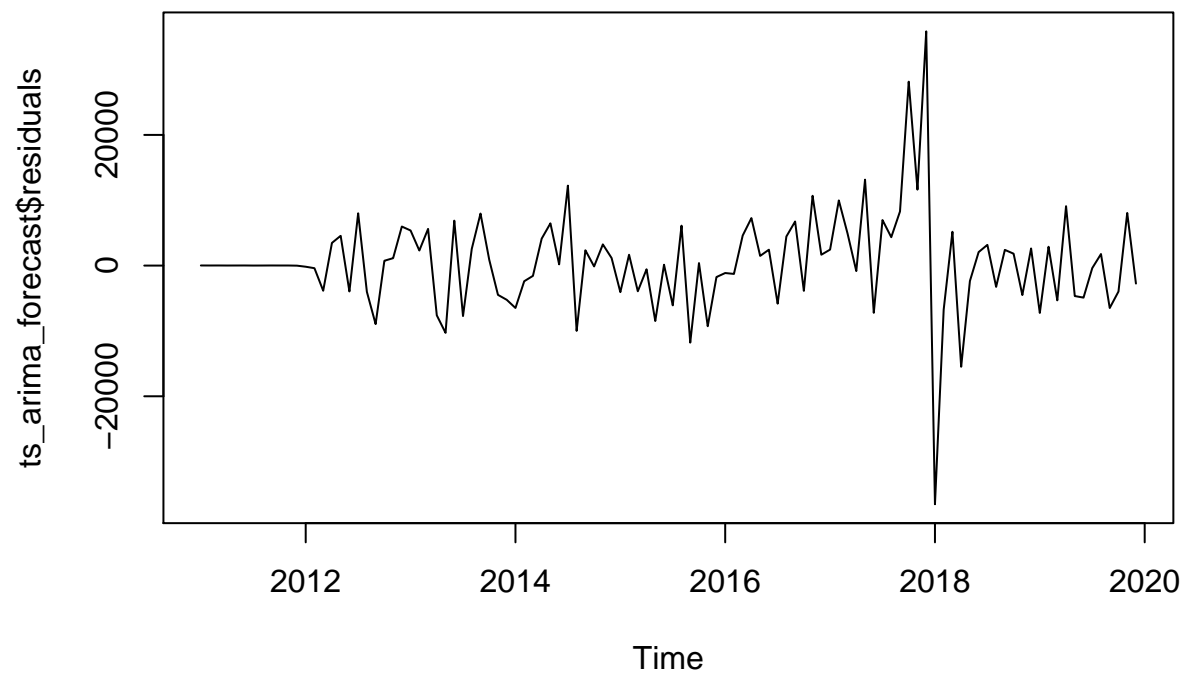


```
Box.test(ts_arma_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: ts_arma_forecast$residuals  
## X-squared = 19.714, df = 20, p-value = 0.4759
```

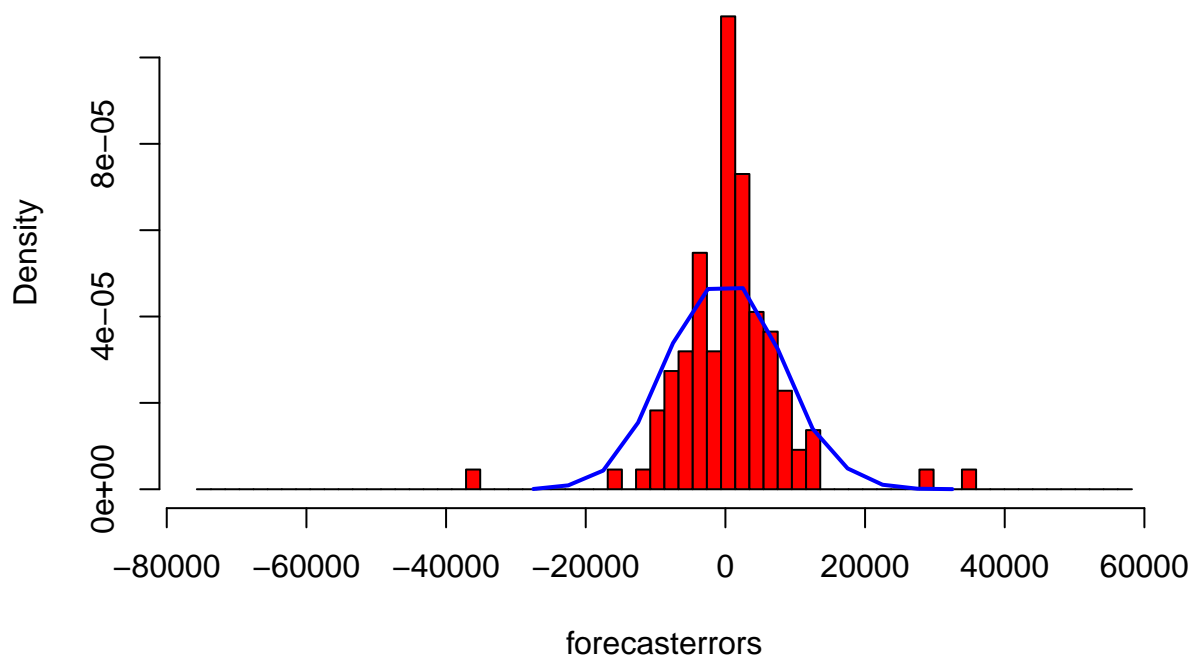
we can reject the null hypothesis, it's rather similar to the HW

```
plot.ts(ts_arma_forecast$residuals)           # make time plot of forecast errors
```



```
plotForecastErrors(ts_arma_forecast$residuals)
```

## Histogram of forecasterrors



```
# Arima, 0,1,0 as given from the loop
```

```
ts_arima = Arima(ts, order=c(2,1,1),seasonal = list(order = c(2,1,0)))
ts_arima
```

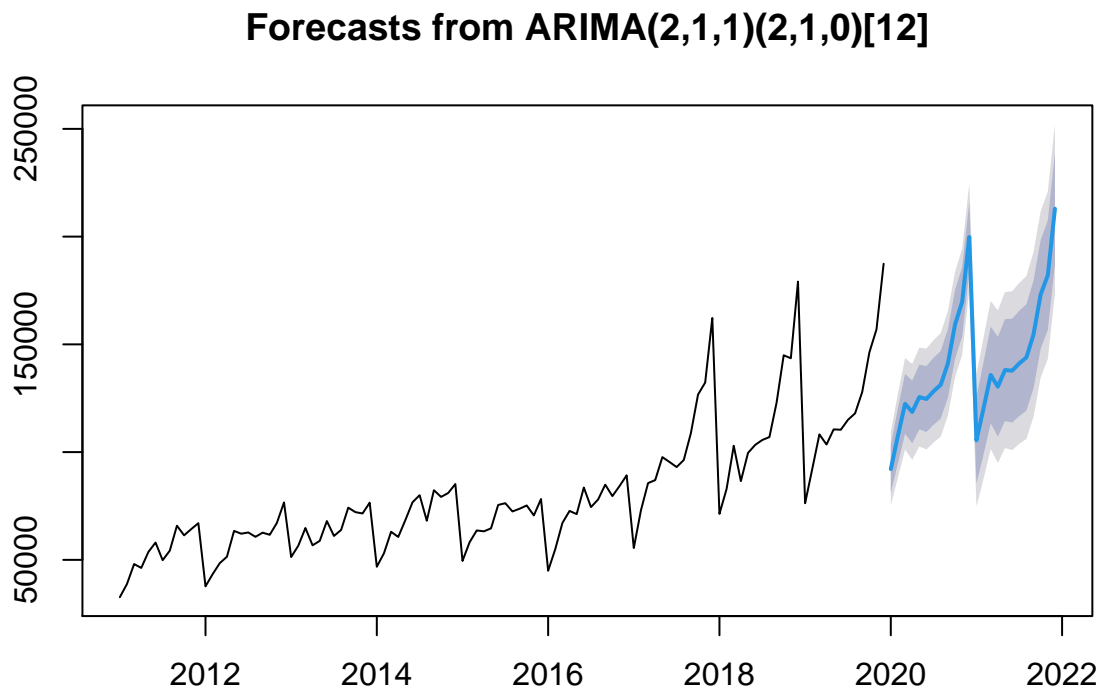
```
## Series: ts
## ARIMA(2,1,1)(2,1,0)[12]
##
## Coefficients:
##          ar1      ar2      ma1      sar1      sar2
##          0.5262  0.1481  -0.9532  -0.0865  -0.1356
## s.e.      0.1144  0.1093   0.0558   0.1043   0.1196
##
## sigma^2 estimated as 74528880:  log likelihood=-994.04
## AIC=2000.08   AICc=2001.03   BIC=2015.4
```

```
ts_arima_forecast = forecast(ts_arima,h = 24)
ts_arima_forecast
```

```
##          Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Jan 2020          92115.67  81051.95 103179.4  75195.18 109036.2
## Feb 2020         107593.74  94842.37 120345.1  88092.20 127095.3
## Mar 2020         122397.56 108513.50 136281.6 101163.71 143631.4
## Apr 2020         118630.34 104081.57 133179.1  96379.91 140880.8
## May 2020         125545.51 110553.33 140537.7 102616.95 148474.1
## Jun 2020         124649.43 109346.86 139952.0 101246.17 148052.7
```

```
## Jul 2020      128240.39 112707.02 143773.8 104484.15 151996.6
## Aug 2020      131240.95 115526.91 146955.0 107208.39 155273.5
## Sep 2020      141128.71 125266.42 156991.0 116869.43 165388.0
## Oct 2020      159176.98 143188.04 175165.9 134724.00 183630.0
## Nov 2020      169717.11 153616.21 185818.0 145092.90 194341.3
## Dec 2020      199814.65 183611.99 216017.3 175034.82 224594.5
## Jan 2021      105460.24  85380.97 125539.5  74751.65 136168.8
## Feb 2021      120373.79  98939.39 141808.2  87592.71 153154.9
## Mar 2021      135808.06 113352.50 158263.6 101465.24 170150.9
## Apr 2021      130363.30 107213.92 153512.7  94959.37 165767.2
## May 2021      138099.58 114428.36 161770.8 101897.58 174301.6
## Jun 2021      137787.89 113706.54 161869.2 100958.64 174617.1
## Jul 2021      141150.28 116730.36 165570.2 103803.24 178497.3
## Aug 2021      143923.52 119212.90 168634.1 106131.88 181715.1
## Sep 2021      154626.96 129658.33 179595.6 116440.73 192813.2
## Oct 2021      173197.54 147993.71 198401.4 134651.61 211743.5
## Nov 2021      182114.41 156691.65 207537.2 143233.65 220995.2
## Dec 2021      212935.48 187305.56 238565.4 173737.90 252133.1
```

```
forecast::plot.forecast(ts_arma_forecast)
```



```
## Growth
```

```
# this_year_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1]
#
# year_2019_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1][c(1:2),]
```

```
# sum_year_2019 = sum(c(year_2019,year_2019_predict_ARIMA))
# year_2020 = (as.data.frame(ts_arima_forecast))[1][c(3:14),]

year_2020 <- (as.data.frame(ts_arima_forecast))[1][c(1:12),]
year_2021 <- (as.data.frame(ts_arima_forecast))[1][c(13:24),]

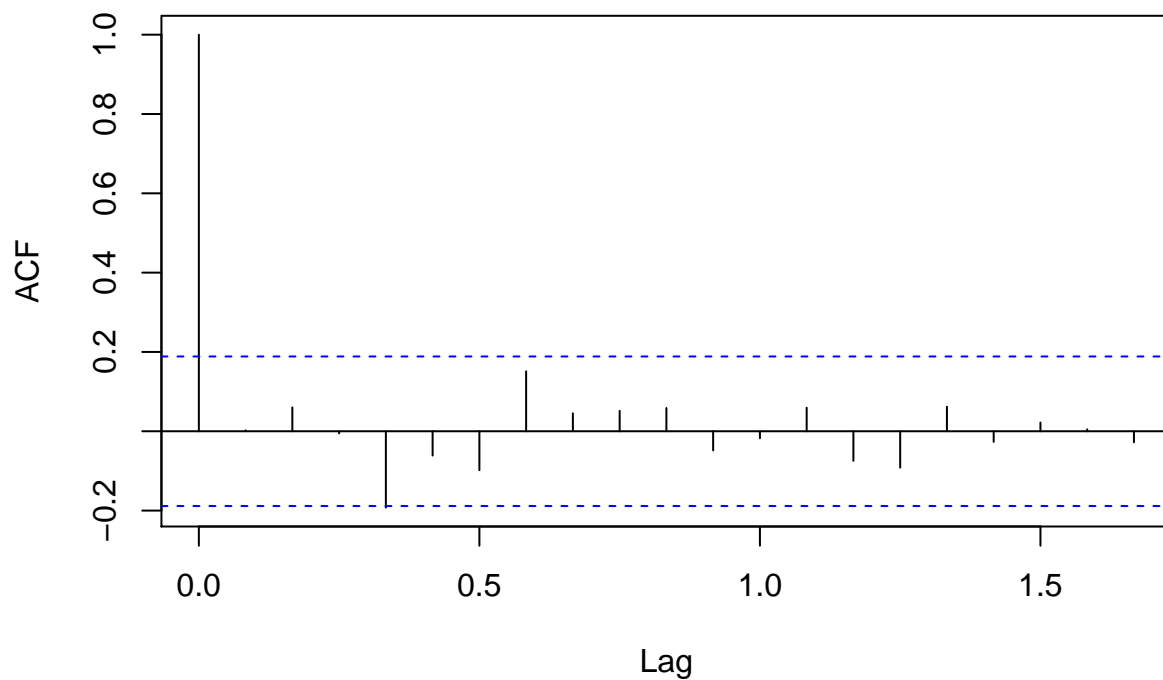
growth_ARIMA2_21 <- growth(sum(year_2021), sum(year_2020))
growth_ARIMA2_20 <- growth(sum(year_2020), sum(year_2019))
```

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arima_forecast$residuals, lag.max=20)
```

### Series ts\_arima\_forecast\$residuals

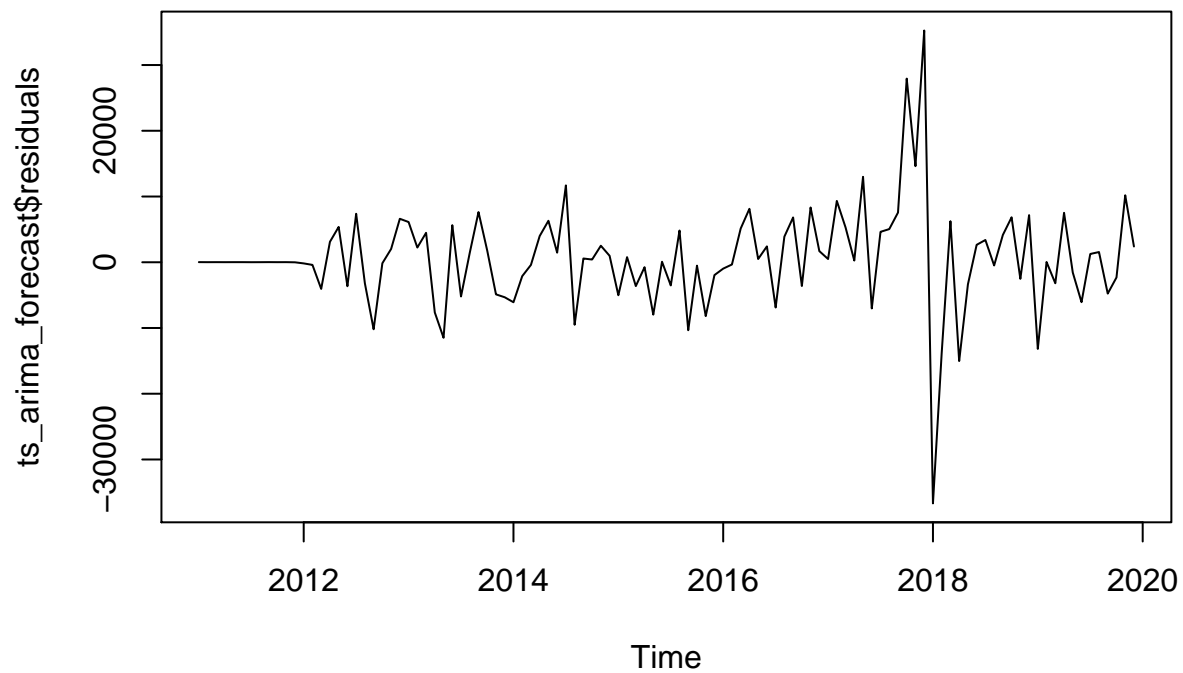


```
Box.test(ts_arima_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data: ts_arima_forecast$residuals
## X-squared = 13.167, df = 20, p-value = 0.8701
```

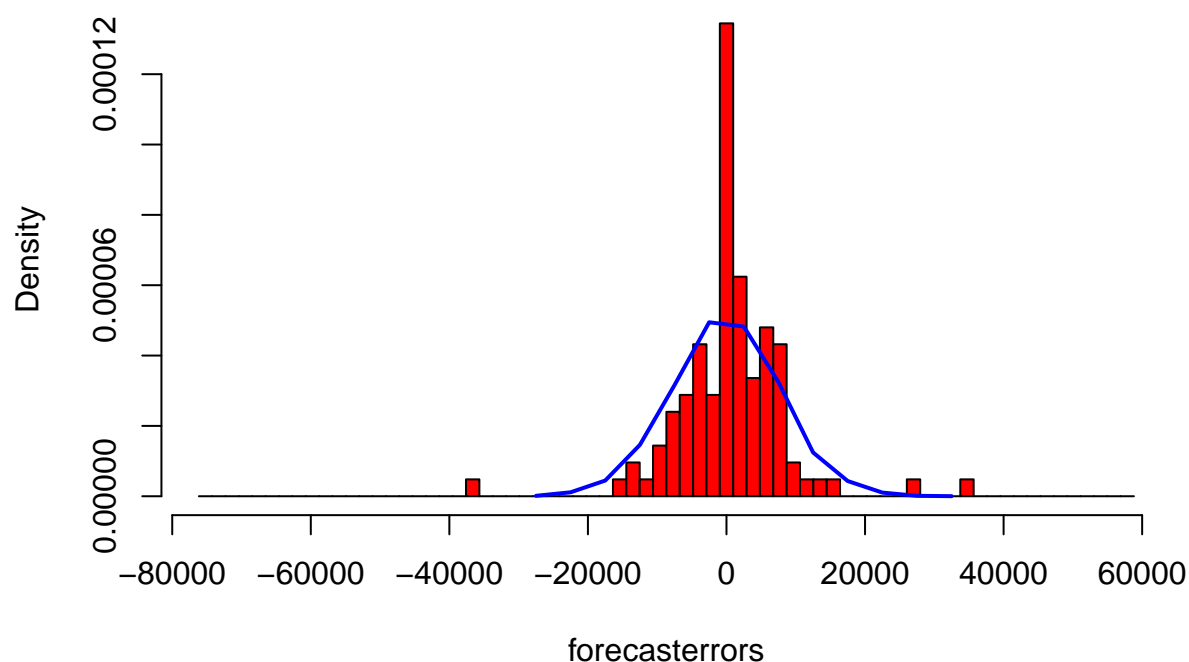
we can reject the null hypothesis, it's rather similar to the HW

```
plot.ts(ts_arima_forecast$residuals)           # make time plot of forecast errors
```



```
plotForecastErrors(ts_arima_forecast$residuals)
```

## Histogram of forecasterrors



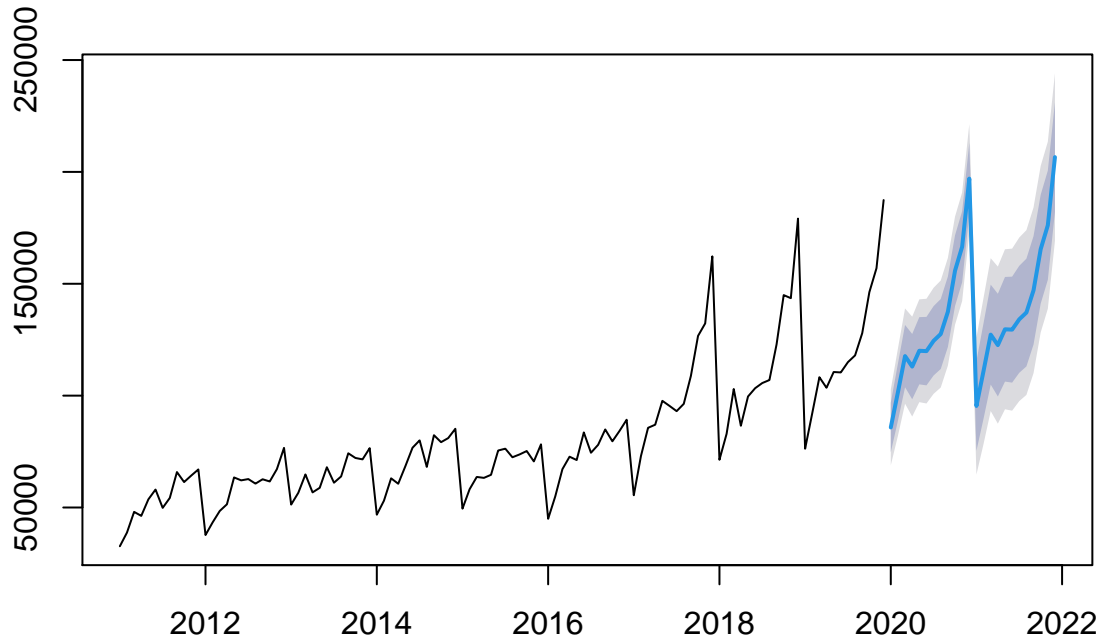
## A model chosen automatically

```
fit <- auto.arima(ts,max.p = 5,max.q = 5,max.P = 5,max.Q = 5,max.d = 3,seasonal = TRUE)
fit
```

```
## Series: ts
## ARIMA(2,1,1)(0,1,0)[12]
##
## Coefficients:
##          ar1      ar2      ma1
##      0.5543  0.1695 -0.9819
## s.e.  0.1080  0.1060  0.0549
##
## sigma^2 estimated as 74285214:  log likelihood=-994.88
## AIC=1997.77   AICc=1998.21   BIC=2007.98
```

```
fit_forecast = forecast(fit,h=24)
plot(fit_forecast)
```

## Forecasts from ARIMA(2,1,1)(0,1,0)[12]



```
# str(fit)
```

## Growth

```
#
# year_2019_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(1:2),]
# year_2019_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(1:2),]
# year_2019_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(1:2),]
#
# sum_year_2019 = sum(c(year_2019,year_2019_predict_auto.arima))
# sum_year_2019_low = sum(c(year_2019,year_2019_predict_auto.arima_95_low))
# sum_year_2019_high = sum(c(year_2019,year_2019_predict_auto.arima_95_high))
#
#
# year_2020_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(3:14),]
# year_2020_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(3:14),]
# year_2020_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(3:14),]
#
#
# growth_auto.arima <- growth(sum(year_2020_predict_auto.arima),sum_year_2019)
# growth_auto.arima_95_low <- growth(sum(year_2020_predict_auto.arima_95_low),sum_year_2019_low)
# growth_auto.arima_95_high <- growth(sum(year_2020_predict_auto.arima_95_high),sum_year_2019_high)
```



```

year_2020 <- (as.data.frame(fit_forecast))[1][c(1:12),]
year_2021 <- (as.data.frame(fit_forecast))[1][c(13:24),]

growth_auto.arima_21 <- growth(sum(year_2021), sum(year_2020))
growth_auto.arima_20 <- growth(sum(year_2020), sum(year_2019))

```

## all the growths

```
growth_ARIMA_20
```

```
## [1] 0.1004553
```

```
growth_ARIMA_21
```

```
## [1] 0.08205767
```

```
growth_ARIMA2_20
```

```
## [1] 0.1154102
```

```
growth_ARIMA2_21
```

```
## [1] 0.09602832
```

```
growth_auto.arima_20
```

```
## [1] 0.0785966
```

```
growth_auto.arima_21
```

```
## [1] 0.07352846
```

```
growth_HW_20
```

```
## [1] 0.03842402
```

```
growth_HW_21
```

```
## [1] 0.03934687
```