

Time Series Forecasting report

Kevork Sulahian

September 26, 2019

```
library(readxl)
library(forecast)
```


```
df <- read_xls('economy.xls', sheet='2011-2019 NACE 2')
```

```
## New names:
## * `` -> ...2
## * `` -> ...3
## * `` -> ...4
## * `` -> ...5
## * `` -> ...6
## * ... and 99 more problems
```

```
df = df[4,]
df = df[-c(1,3)]
rownames(df) = df[1]
```


```
## Warning: Setting row names on a tibble is deprecated.
```

```
df = df[-1]
df = t(df)
df[] <- sapply(df[],function(x) as.numeric(as.character(x)))
df = as.numeric(df)
df= df * 1000000
ts = ts(df, start = c(2011,1), frequency = c(12))
```



In this case, it appears that an additive model is not appropriate for describing this time series, since the size of the seasonal fluctuations and random fluctuations seem to increase with the level of the time series. Thus, we may need to transform the time series in order to get a transformed time series that can be described using an additive model. For example, we can transform the time series by calculating the natural log of the original data:

```
log_ts <- log(ts)
plot.ts(log_ts)
```



Decomposing Time Series

Decomposing a time series means separating it into its constituent components, which are usually a trend component and an irregular component, and if it is a seasonal time series, a seasonal component.

Decomposing Seasonal Data

A seasonal time series consists of a trend component, a seasonal component and an irregular component. Decomposing the time series means separating the time series into these three components: that is, estimating these three components.

```
ts_components <- decompose(ts)
```

we can print out the estimated values of the seasonal component

```
ts_components$seasonal
```


```
##           Jan           Feb           Mar           Apr           May
## 2011 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
## 2012 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
## 2013 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
## 2014 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
## 2015 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
```

```

## 2016 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
## 2017 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
## 2018 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
## 2019 -18114721416 -13972761892 -3652881535 -10968033321 -5840660702
##
##           Jun           Jul           Aug           Sep           Oct
## 2011 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2012 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2013 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2014 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2015 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2016 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2017 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2018 -1030591654 -3088941580 -2737647830  5107804253 11091117274
## 2019 -1030591654
##
##           Nov           Dec
## 2011 13028793837 30178524566
## 2012 13028793837 30178524566
## 2013 13028793837 30178524566
## 2014 13028793837 30178524566
## 2015 13028793837 30178524566
## 2016 13028793837 30178524566
## 2017 13028793837 30178524566
## 2018 13028793837 30178524566
## 2019

```


Decomposition of additive time series



The plot above shows the original time series (top), the estimated trend component (second from top), the estimated seasonal component (third from top), and the estimated irregular component (bottom)

Seasonally Adjusting

```
ts_seasonall <- ts - ts_components$seasonal
```



```
## Holt-Winters Exponential Smoothing
```

```
ts_forcaste <- HoltWinters(ts)
ts_forcaste
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = ts)
##
## Smoothing parameters:
##  alpha: 0.3841348
##  beta : 0
##  gamma: 1
##
## Coefficients:
##           [,1]
## a  166929650345
## b    859005492
## s1  2680812331
## s2  1004407035
```

```
## s3      9728353964
## s4      31427608853
## s5      23274665015
## s6      56266579294
## s7     -47717485979
## s8     -33061264111
## s9      -7294279687
## s10    -17388238299
## s11     -7579809787
## s12      458749655
```


```
#
```

The value of α (0.41) is relatively low, indicating that the estimate of the level at the current time point is based upon both recent observations and some observations in the more distant past. The value of β is 0.00, indicating that the estimate of the slope b of the trend component is not updated over the time series, and instead is set equal to its initial value. This makes good intuitive sense, as the level changes quite a bit over the time series, but the slope b of the trend component remains roughly the same. In contrast, the value of γ (0.96) is high, indicating that the estimate of the seasonal component at the current time point is just based upon very recent observations

```
ts_forcaste$SSE
```

```
## [1] 9.131448e+21
```


Holt-Winters filtering



```
ts_forcaste2 = forecast::forecast.HoltWinters(ts_forcaste, h= 18)
(as.data.frame(ts_forcaste2))[1]
```

```
##          Point Forecast
## Jul 2019   170469468169
## Aug 2019   169652068366
## Sep 2019   179235020787
## Oct 2019   201793281168
## Nov 2019   194499342823
## Dec 2019   228350262594
## Jan 2020   125225202813
## Feb 2020   140740430174
## Mar 2020   167366420090
## Apr 2020   158131466971
## May 2020   168798900975
## Jun 2020   177696465909
## Jul 2020   180777534078
## Aug 2020   179960134275
## Sep 2020   189543086696
## Oct 2020   212101347077
## Nov 2020   204807408732
## Dec 2020   238658328503
```

Forecasts from HoltWinters




Growth

```
year_2019 <- window(ts, 2019)
year_2019_predict_HW <- (as.data.frame(ts_forcaste2))[1][c(1:6),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_HW))
year_2020 = (as.data.frame(ts_forcaste2))[1][c(7:18),]
growth_HW <- growth(sum(year_2020),sum_year_2019)
growth_HW
```

```
## [1] 0.0485728
```

We can investigate whether the predictive model can be improved upon by checking whether the in-sample forecast errors show non-zero autocorrelations at lags 1-20, by making a correlogram and carrying out the

Series ts_forcaste2\$residuals




Ljung-Box test:

```
##
## Box-Ljung test
##
## data: ts_forcaste2$residuals
## X-squared = 23.626, df = 20, p-value = 0.2591
```


The correlogram shows that the autocorrelations for the in-sample forecast errors do not exceed the significance bounds for lags 1-20. Furthermore, the p-value for Ljung-Box test is 0.2, indicating that there is little evidence of non-zero autocorrelations at lags 1-20.

We can check whether the forecast errors have constant variance over time, and are normally distributed with mean zero, by making a time plot of the forecast errors and a histogram (with overlaid normal curve):


```
plot.ts(ts_forcaste2$residuals)
```




```
plotForecastErrors(ts_forcaste2$residuals)
```




From the time plot, it appears plausible that the forecast errors have constant variance over time. From the histogram of forecast errors, it seems plausible that the forecast errors are normally distributed with mean zero.

Thus, there is little evidence of autocorrelation at lags 1-20 for the forecast errors, and the forecast errors appear to be normally distributed with mean zero and constant variance over time. This suggests that Holt-Winters exponential smoothing provides an adequate predictive model of the log of total productivity, which probably cannot be improved upon. Furthermore, the assumptions upon which the prediction intervals were based are probably valid.

```
plot.ts(ts)
```




```
ts_diff1 <- diff(ts, differences = 1)
plot.ts(ts_diff1)
```



The time series of differences (above) does appear to be stationary in mean and variance, as the level of the series stays roughly constant over time, and the variance of the series appears roughly constant over time

```
acf(ts_diff1, lag.max=20) # plot a correlogram
```




We see from the correlogram that the autocorrelation exceeds the significance bound 3 times but all the others do not exceed

```
acf(ts_diff1, lag.max=20, plot=FALSE) # get the autocorrelation values
```

```
##
## Autocorrelations of series 'ts_diff1', by lag
##
## 0.0000 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500
## 1.000 -0.342 -0.101 0.101 -0.180 -0.014 0.051 0.015 -0.145 0.115
## 0.8333 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833
## -0.083 -0.228 0.644 -0.199 -0.138 0.112 -0.116 -0.044 0.095 -0.007
## 1.6667
## -0.090
```

```
pacf(ts_diff1, lag.max=20) # plot a partial correlogram
```



```
pacf(ts_diff1, lag.max=20, plot=FALSE) # get the partial autocorrelation values
```

```
##
## Partial autocorrelations of series 'ts_diff1', by lag
##
## 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333
## -0.342 -0.246 -0.031 -0.209 -0.183 -0.117 -0.041 -0.244 -0.108 -0.219
## 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
## -0.525 0.315 0.129 -0.070 -0.056 0.044 0.009 0.071 0.008 0.064
```

Arima, 1,1,1

```
ts_arima = Arima(ts, order=c(1,1,1),seasonal = list(order = c(1,1,1)))
ts_arima
```

```
## Series: ts
## ARIMA(1,1,1)(1,1,1)[12]
##
## Coefficients:
##          ar1      ma1      sar1      sma1
##      -0.3595  0.0227 -0.5334  0.4228
## s.e.   0.3126  0.3363  1.1785  1.2365
```


```
##
## sigma^2 estimated as 1.087e+20: log likelihood=-2177.43
## AIC=4364.86 AICc=4365.58 BIC=4377.3
```

```
ts_arima_forecast = forecast(ts_arima,h = 18)
ts_arima_forecast
```

##		Point Forecast	Lo 80	Hi 80	Lo 95
##	Jul 2019	170788319965	157428067448	184148572483	150355576982
##	Aug 2019	170645005267	154613295424	186676715109	146126620875
##	Sep 2019	180021974811	160869582624	199174366999	150730918630
##	Oct 2019	210457121225	188898689473	232015552977	177486345071
##	Nov 2019	206924998308	183115365432	230734631185	170511307141
##	Dec 2019	251110539040	225274163882	276946914198	211597213035
##	Jan 2020	141729506303	114004417826	169454594780	99327642544
##	Feb 2020	151566190295	122076402987	181055977603	106465452572
##	Mar 2020	175554800329	144398986497	206710614162	127906094914
##	Apr 2020	164410562793	131673783887	197147341700	114343979948
##	May 2020	172509839910	138264873602	206754806217	120136683480
##	Jun 2020	181614978374	145925547412	217304409335	127032703892
##	Jul 2020	184727206520	142836852777	226617560262	120661439022
##	Aug 2020	184216559042	138649586711	229783531373	114527888641
##	Sep 2020	193265167425	143752001790	242778333059	117541315171
##	Oct 2020	223788192073	170807130446	276769253701	142760650580
##	Nov 2020	219721294255	163423195686	276019392823	133620782532
##	Dec 2020	263567611400	204158698957	322976523843	172709520411
##		Hi 95			
##	Jul 2019	191221062948			
##	Aug 2019	195163389659			
##	Sep 2019	209313030992			
##	Oct 2019	243427897379			
##	Nov 2019	243338689476			
##	Dec 2019	290623865044			
##	Jan 2020	184131370062			
##	Feb 2020	196666928018			
##	Mar 2020	223203505744			
##	Apr 2020	214477145638			
##	May 2020	224882996339			
##	Jun 2020	236197252855			
##	Jul 2020	248792974017			
##	Aug 2020	253905229443			
##	Sep 2020	268989019678			
##	Oct 2020	304815733566			
##	Nov 2020	305821805978			
##	Dec 2020	354425702389			

```
forecast::plot.forecast(ts_arima_forecast)
```

Forecasts from ARIMA(1,1,1)(1,1,1)[12]



Growth

```
this_year_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1]

# growth_ARIMA <- growth(sum(c(this_year,as.numeric(this_year_predict_ARIMA$`Point Forecast`))), sum(la
# growth_ARIMA

year_2019_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1][c(1:6),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_ARIMA))
year_2020 = (as.data.frame(ts_arma_forecast))[1][c(7:18),]
growth_ARIMA <- growth(sum_year_2019, sum(year_2020))
-growth_ARIMA
```


[1] 0.0736588

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arma_forecast$residuals, lag.max=20)
```


Series ts_arma_forecast\$residuals




```
Box.test(ts_arma_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: ts_arma_forecast$residuals  
## X-squared = 23.833, df = 20, p-value = 0.2498
```


we can reject the null hypothesis, it's rather similar to the HW

```
plot.ts(ts_arma_forecast$residuals)           # make time plot of forecast errors
```



```
plotForecastErrors(ts_arma_forecast$residuals)
```

Histogram of forecasterrors



Arima, 0,1,0 as given from the loop

```
ts_arima = Arima(ts, order=c(2,1,1),seasonal = list(order = c(2,1,0)))
ts_arima
```

```
## Series: ts
## ARIMA(2,1,1)(2,1,0)[12]
##
## Coefficients:
##          ar1      ar2      ma1      sar1      sar2
##          0.5222  0.1982 -0.9857 -0.1363 -0.1190
## s.e.      0.1135  0.1114  0.0861  0.1112  0.1648
##
## sigma^2 estimated as 9.836e+19: log likelihood=-2173.54
## AIC=4359.08   AICc=4360.11   BIC=4374.01
```

```
ts_arima_forecast = forecast(ts_arima,h = 18)
ts_arima_forecast
```

```
##          Point Forecast      Lo 80      Hi 80      Lo 95
## Jul 2019  168553527898 155828774248 181278281548 149092696561
## Aug 2019  167531660835 153077999389 181985322282 145426693199
## Sep 2019  175847002176 160075741268 191618263083 151726940057
## Oct 2019  201591574585 185088058560 218095090611 176351624904
## Nov 2019  198603614313 181623400634 215583827992 172634618484
## Dec 2019  237810566298 220521444450 255099688146 211369136193
```


```

## Jan 2020 138627062656 121129531745 156124593567 111866898395
## Feb 2020 150504321916 132862690003 168145953828 123523774190
## Mar 2020 173220725922 155476469987 190964981857 146083228298
## Apr 2020 163632514180 145813038373 181451989987 136379977689
## May 2020 172515909415 154639634745 190392184085 145176506559
## Jun 2020 178798587286 160878126902 196719047669 151391608213
## Jul 2020 180344264986 158682056138 202006473835 147214775460
## Aug 2020 179825794236 157062135945 202589452527 145011783153
## Sep 2020 188730868045 165072010340 212389725751 152547767603
## Oct 2020 213672268657 189463832448 237880704866 176648660426
## Nov 2020 211845497910 187253149128 236437846692 174234746065
## Dec 2020 250844476239 225983538992 275705413486 212822953784
## Hi 95
## Jul 2019 188014359235
## Aug 2019 189636628471
## Sep 2019 199967064295
## Oct 2019 226831524266
## Nov 2019 224572610142
## Dec 2019 264251996403
## Jan 2020 165387226917
## Feb 2020 177484869641
## Mar 2020 200358223546
## Apr 2020 190885050671
## May 2020 199855312271
## Jun 2020 206205566359
## Jul 2020 213473754513
## Aug 2020 214639805319
## Sep 2020 224913968488
## Oct 2020 250695876889
## Nov 2020 249456249755
## Dec 2020 288865998693

```

```
forecast:::plot.forecast(ts_arima_forecast)
```

Forecasts from ARIMA(2,1,1)(2,1,0)[12]



Growth

```
this_year_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1]


year_2019_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1][c(1:6),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_ARIMA))
year_2020 = (as.data.frame(ts_arma_forecast))[1][c(7:18),]
growth_ARIMA2 <- growth(sum_year_2019, sum(year_2020))
-growth_ARIMA2
```

[1] 0.0690669

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arima_forecast$residuals, lag.max=20)
```




```
Box.test(ts_arima_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: ts_arima_forecast$residuals  
## X-squared = 17.07, df = 20, p-value = 0.6484
```

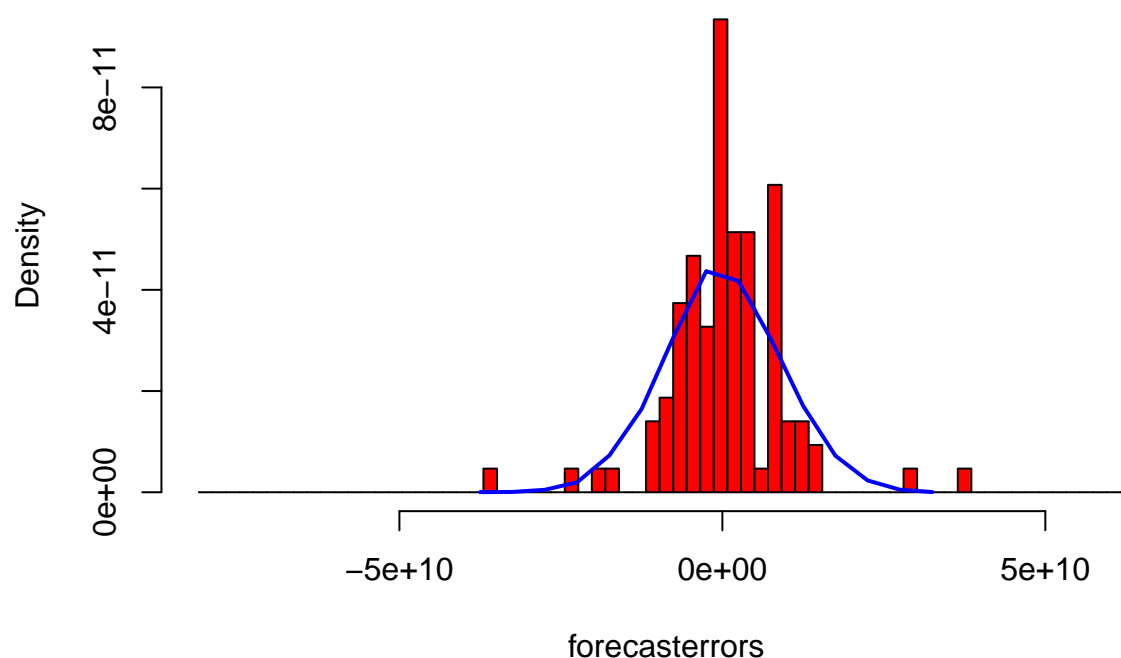
we can reject the null hypothesis, it's rather similar to the HW

```
plot.ts(ts_arima_forecast$residuals)           # make time plot of forecast errors
```



```
plotForecastErrors(ts_arma_forecast$residuals)
```

Histogram of forecasterrors




A model chosen automatically

```
fit <- auto.arima(ts,max.p = 5,max.q = 5,max.P = 5,max.Q = 5,max.d = 3,seasonal = TRUE)
fit
```

```
## Series: ts
## ARIMA(0,0,0)(0,1,0)[12] with drift
##
## Coefficients:
##          drift
##      924981389
## s.e.   84479299
##
## sigma^2 estimated as 1.687e+20:  log likelihood=-2223.06
## AIC=4450.13   AICc=4450.27   BIC=4455.13
```

```
fit_forecast = forecast(fit,h=18)
plot(fit_forecast)
```


Forecasts from ARIMA(0,0,0)(0,1,0)[12] with drift



```
# str(fit)
```

Growth

```
year_2019_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(1:6),]
year_2019_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(1:6),]
year_2019_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(1:6),]

sum_year_2019 = sum(c(year_2019,year_2019_predict_auto.arima))
sum_year_2019_low = sum(c(year_2019,year_2019_predict_auto.arima_95_low))
sum_year_2019_high = sum(c(year_2019,year_2019_predict_auto.arima_95_high))

year_2020_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(7:18),]
year_2020_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(7:18),]
year_2020_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(7:18),]

growth_auto.arima <- growth(sum(year_2020_predict_auto.arima),sum_year_2019)
growth_auto.arima_95_low <- growth(sum(year_2020_predict_auto.arima_95_low),sum_year_2019_low)
growth_auto.arima_95_high <- growth(sum(year_2020_predict_auto.arima_95_high),sum_year_2019_high)

growth_auto.arima
```

```
## [1] 0.06473158
```

```
growth_auto.arima_95_low
```

```
## [1] -0.04347715
```

```
growth_auto.arima_95_high
```

```
## [1] 0.1579852
```

all the growths

```
growth_ARIMA
```

```
## [1] -0.0736588
```

```
growth_ARIMA2
```

```
## [1] -0.0690669
```

```
growth_auto.arima
```

```
## [1] 0.06473158
```

```
growth_HW
```

```
## [1] 0.0485728
```