

Time Series Forecasting report for total industry

Kevork Sulahian

2021-04-22

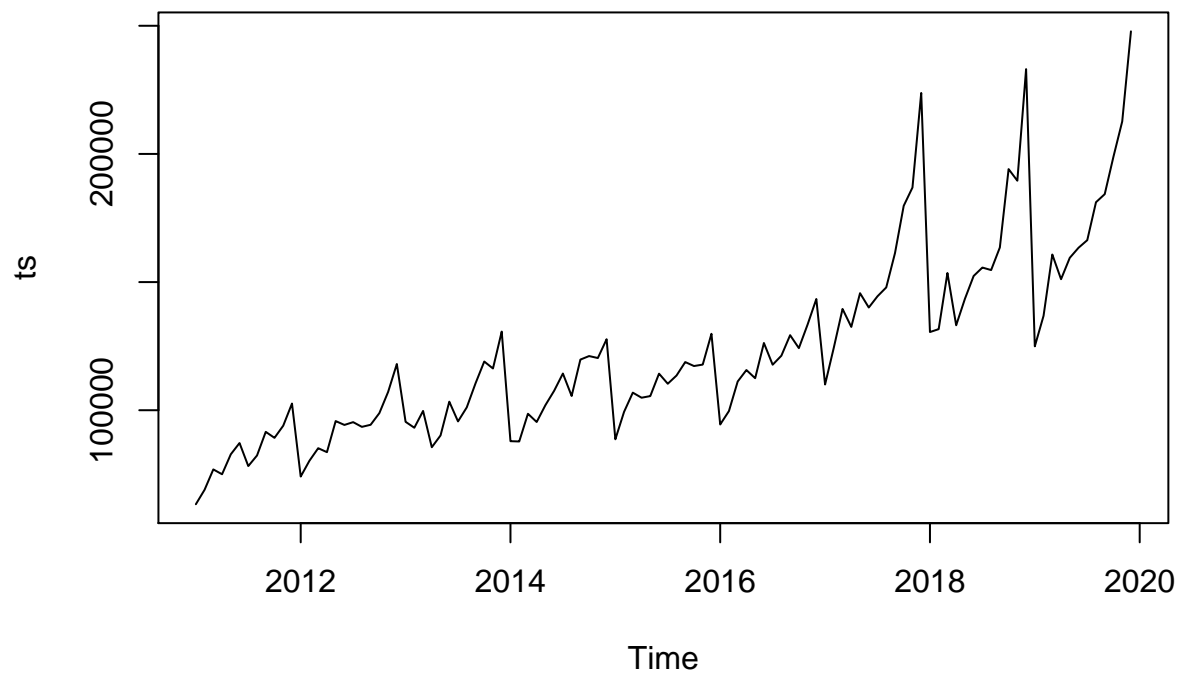
```
library(readxl)
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

df <- read_xls('economy.xls', sheet='2011-2019 NACE 2')

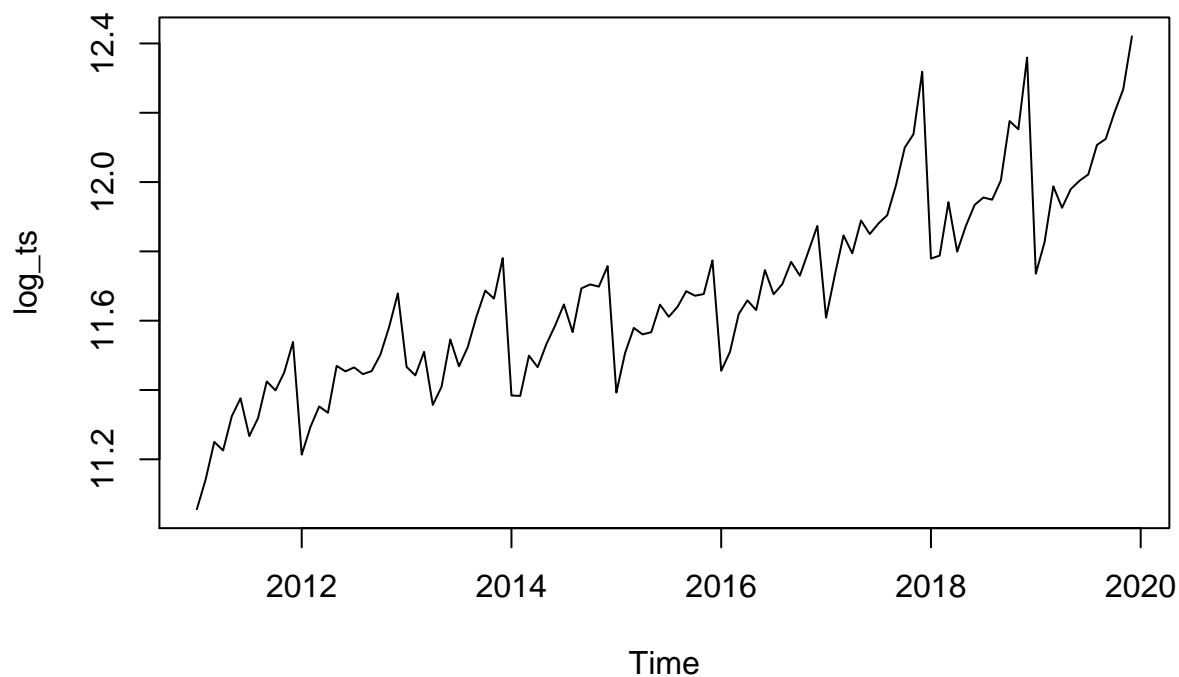
## New names:
## * ' -> ...2
## * ' -> ...3
## * ' -> ...4
## * ' -> ...5
## * ' -> ...6
## * ...

df = df[4,]
df = df[-c(1,3)]
# rownames(df) = df[1]
df = df[-1]
df = t(df)
df[110] = "155770.7"
df[] <- sapply(df[],function(x) as.numeric(as.character(x)))
df = as.numeric(df)
# df= df * 1000000
df = df[-c(109:120)]
ts = ts(df, start = c(2011,1), frequency = c(12))
```



In this case, it appears that an additive model is not appropriate for describing this time series, since the size of the seasonal fluctuations and random fluctuations seem to increase with the level of the time series. Thus, we may need to transform the time series in order to get a transformed time series that can be described using an additive model. For example, we can transform the time series by calculating the natural log of the original data:

```
log_ts <- log(ts)
plot.ts(log_ts)
```



##Decomposing Time Series

Decomposing a time series means separating it into its constituent components, which are usually a trend component and an irregular component, and if it is a seasonal time series, a seasonal component.

###Decomposing Seasonal Data A seasonal time series consists of a trend component, a seasonal component and an irregular component. Decomposing the time series means separating the time series into these three components: that is, estimating these three components.

```
ts_components <- decompose(ts)
```

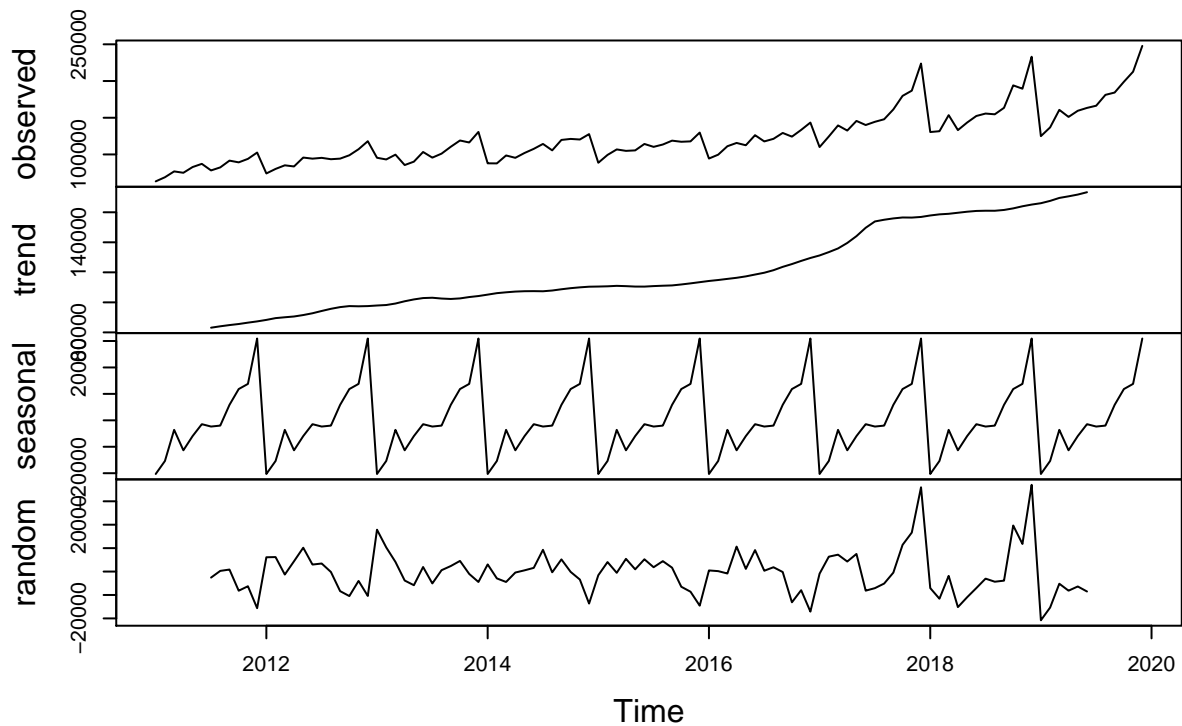
we can print out the estimated values of the seasonal component

```
ts_components$seasonal
```

```
##           Jan           Feb           Mar           Apr           May           Jun
## 2011 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2012 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2013 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2014 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2015 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2016 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2017 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2018 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
## 2019 -20296.500 -15377.583 -3605.059 -11343.268 -5960.194 -1457.614
##           Jul           Aug           Sep           Oct           Nov           Dec
```

```
## 2011 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2012 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2013 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2014 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2015 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2016 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2017 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2018 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
## 2019 -2348.985 -1997.691 5847.761 11831.074 13768.751 30939.306
```

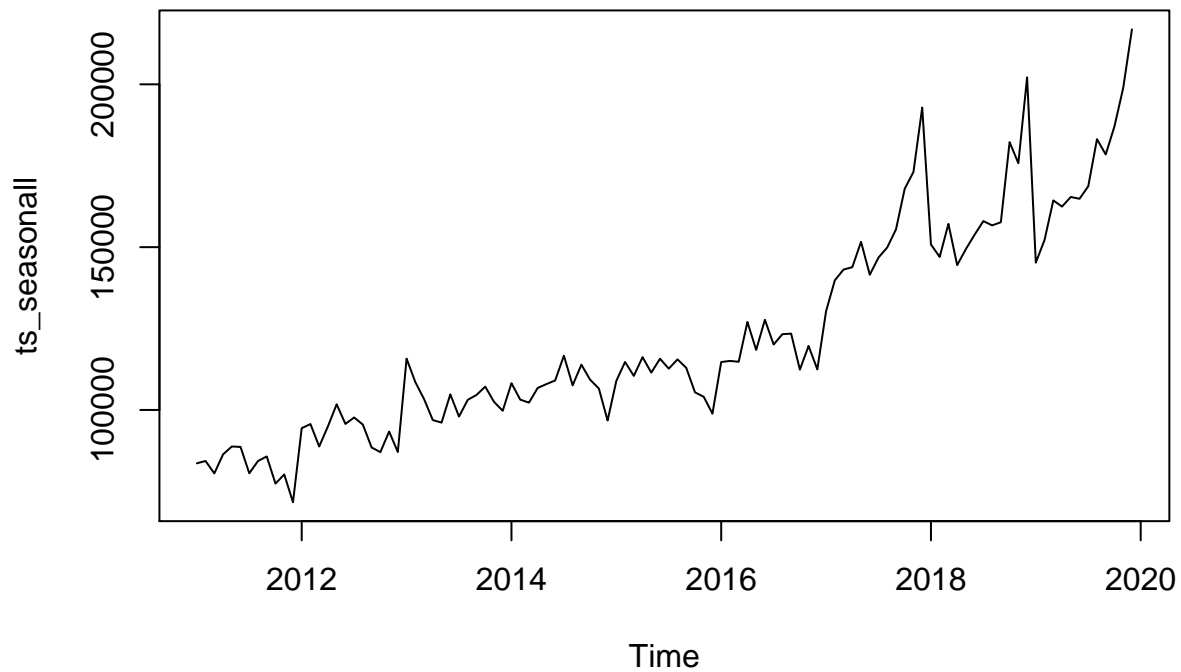
Decomposition of additive time series



The plot above shows the original time series (top), the estimated trend component (second from top), the estimated seasonal component (third from top), and the estimated irregular component (bottom)

Seasonally Adjusting

```
ts_seasonall <- ts - ts_components$seasonal
```



```
## Holt-Winters Exponential Smoothing
```

```
ts_forcaste <- HoltWinters(ts)
ts_forcaste
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = ts)
##
## Smoothing parameters:
##   alpha: 0.3444881
##   beta : 0
##   gamma: 1
##
## Coefficients:
##           [,1]
## a    182005.7804
## b      859.0055
## s1  -46424.5490
## s2  -32968.9766
## s3   -7812.7574
## s4  -18184.3629
## s5   -8636.7940
## s6   -3044.6551
## s7    340.6217
```

```
## s8      9742.9149
## s9     11264.5693
## s10    27672.3528
## s11    35116.2729
## s12    65820.9196
```

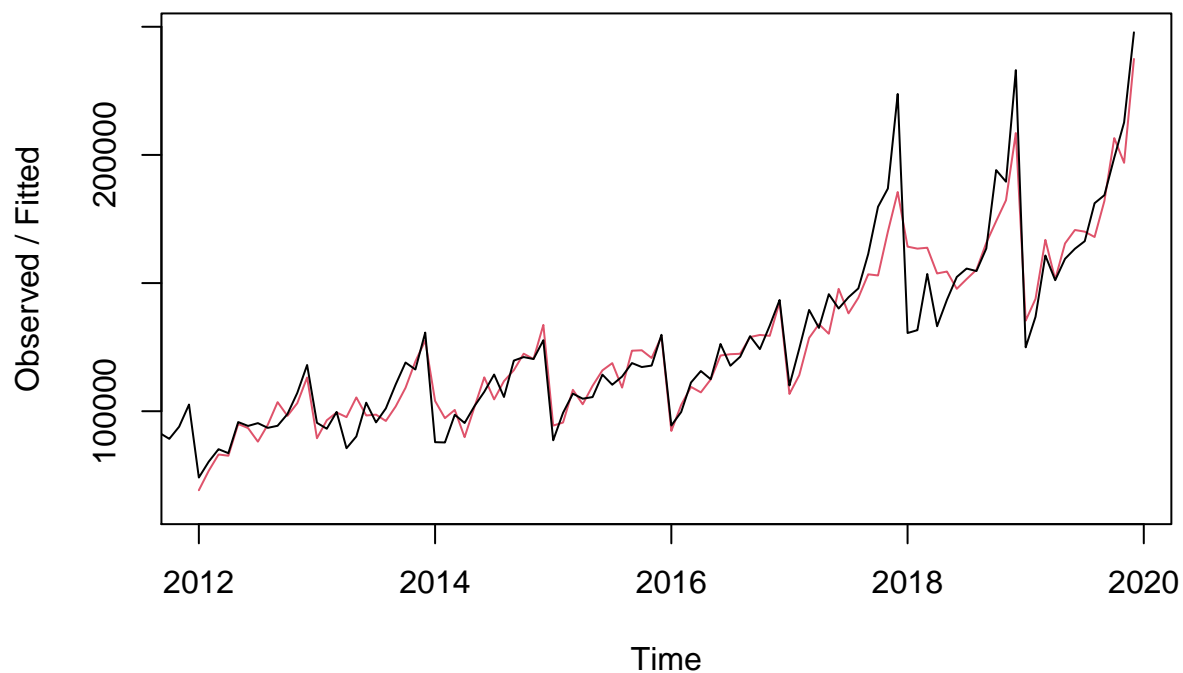
```
#
```

The value of alpha (0.41) is relatively low, indicating that the estimate of the level at the current time point is based upon both recent observations and some observations in the more distant past. The value of beta is 0.00, indicating that the estimate of the slope b of the trend component is not updated over the time series, and instead is set equal to its initial value. This makes good intuitive sense, as the level changes quite a bit over the time series, but the slope b of the trend component remains roughly the same. In contrast, the value of gamma (0.96) is high, indicating that the estimate of the seasonal component at the current time point is just based upon very recent observations

```
ts_forcaste$SSE
```

```
## [1] 9814489310
```

Holt-Winters filtering

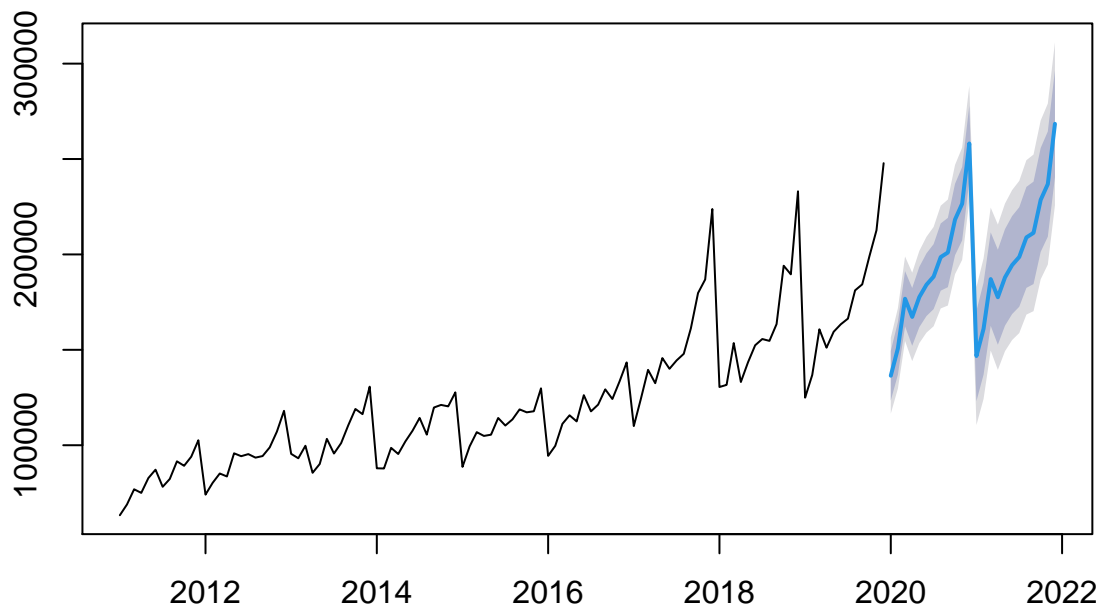


```
ts_forcaste2 = forecast::forecast.HoltWinters(ts_forcaste, h= 24)
(as.data.frame(ts_forcaste2))[1]
```

```
##          Point Forecast
```

## Jan 2020	136440.2
## Feb 2020	150754.8
## Mar 2020	176770.0
## Apr 2020	167257.4
## May 2020	177664.0
## Jun 2020	184115.2
## Jul 2020	188359.4
## Aug 2020	198620.7
## Sep 2020	201001.4
## Oct 2020	218268.2
## Nov 2020	226571.1
## Dec 2020	258134.8
## Jan 2021	146748.3
## Feb 2021	161062.9
## Mar 2021	187078.1
## Apr 2021	177565.5
## May 2021	187972.1
## Jun 2021	194423.2
## Jul 2021	198667.5
## Aug 2021	208928.8
## Sep 2021	211309.5
## Oct 2021	228576.3
## Nov 2021	236879.2
## Dec 2021	268442.8

Forecasts from HoltWinters



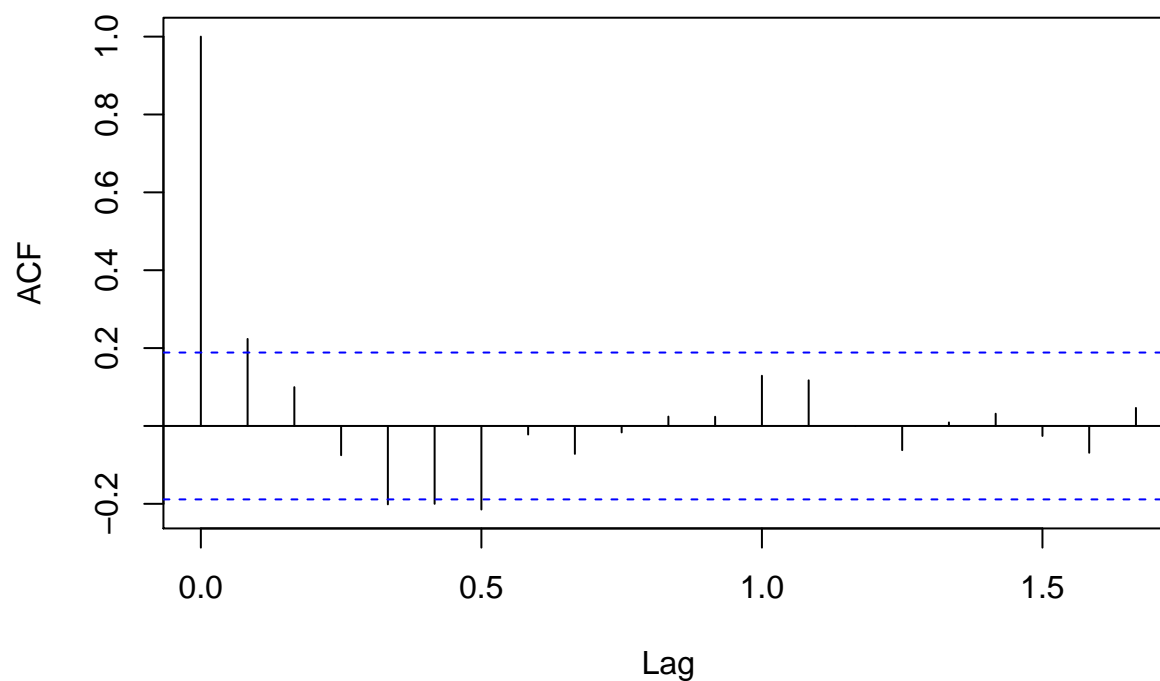
Growth

```
year_2019 <- window(ts, 2019)
year_2020 <- (as.data.frame(ts_forcaste2))[1][c(1:12),]
year_2021 <- (as.data.frame(ts_forcaste2))[1][c(13:24),]

growth_HW_21 <- growth(sum(year_2021),sum(year_2020))
growth_HW_20 <- growth(sum(year_2020),sum(year_2019))
```

We can investigate whether the predictive model can be improved upon by checking whether the in-sample forecast errors show non-zero autocorrelations at lags 1-20, by making a correlogram and carrying out the Ljung-Box test:

Series ts_forcaste2\$residuals

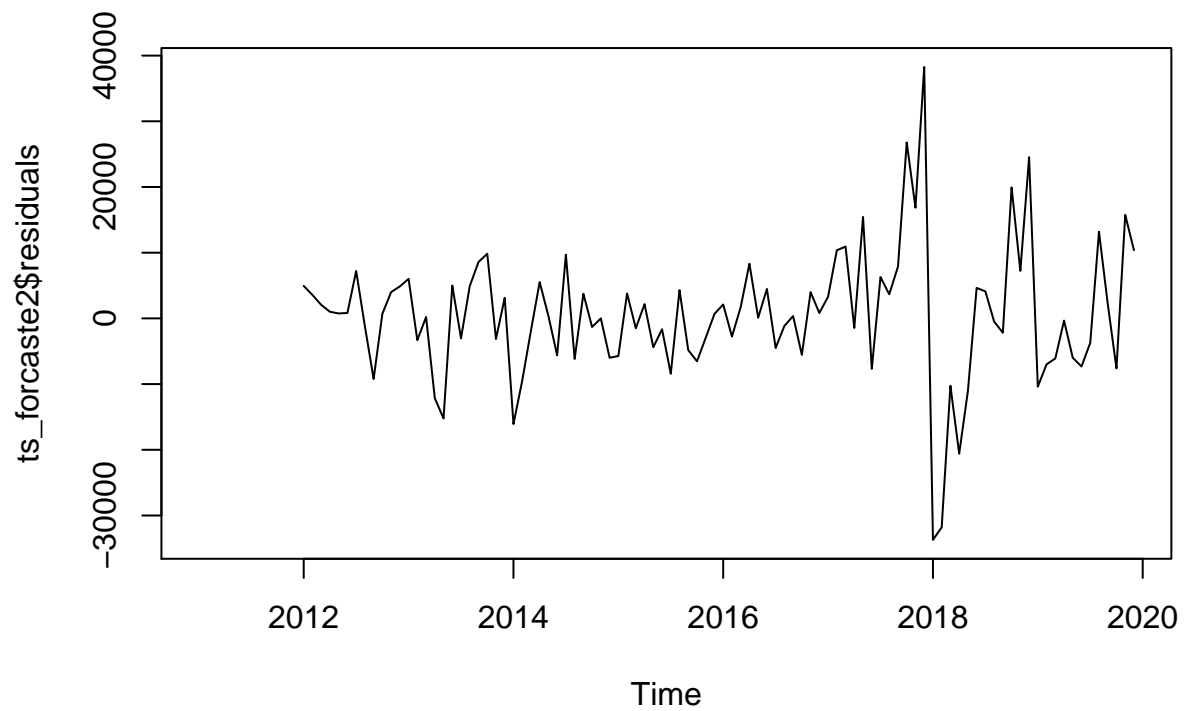


```
##
## Box-Ljung test
##
## data: ts_forcaste2$residuals
## X-squared = 25.284, df = 20, p-value = 0.1908
```

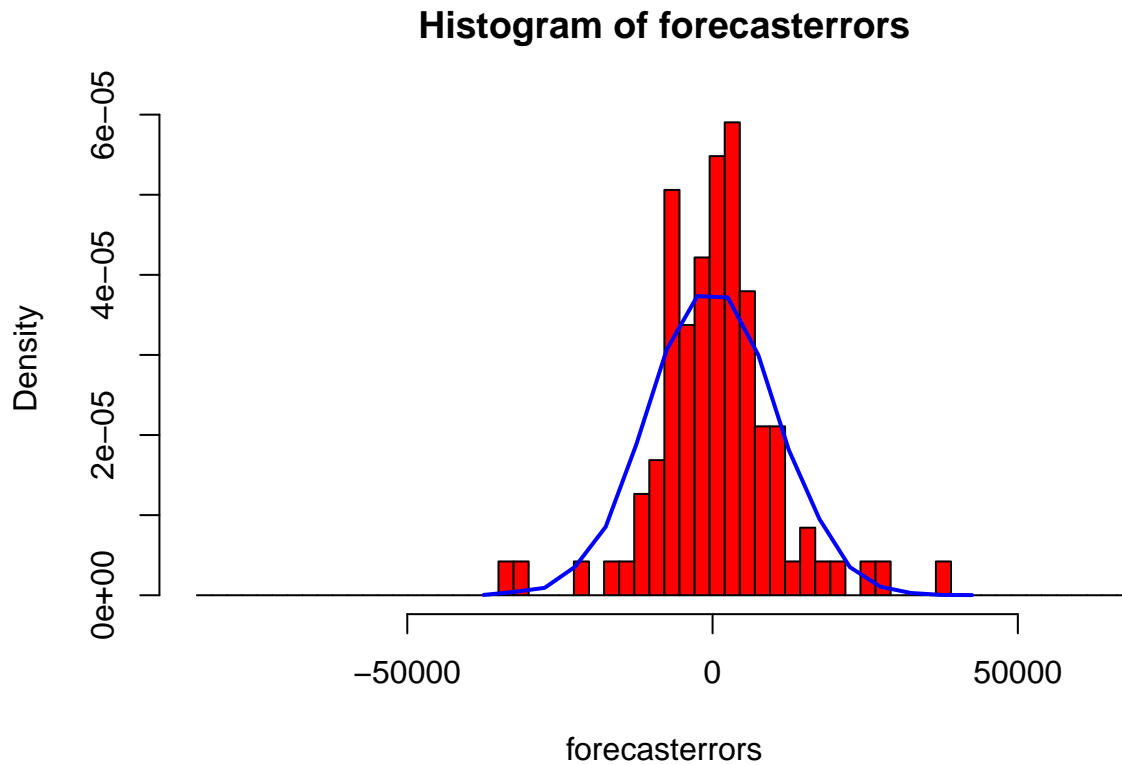
The correlogram shows that the autocorrelations for the in-sample forecast errors do not exceed the significance bounds for lags 1-20. Furthermore, the p-value for Ljung-Box test is 0.2, indicating that there is little evidence of non-zero autocorrelations at lags 1-20.

We can check whether the forecast errors have constant variance over time, and are normally distributed with mean zero, by making a time plot of the forecast errors and a histogram (with overlaid normal curve):


```
plot.ts(ts_forcaste2$residuals)
```



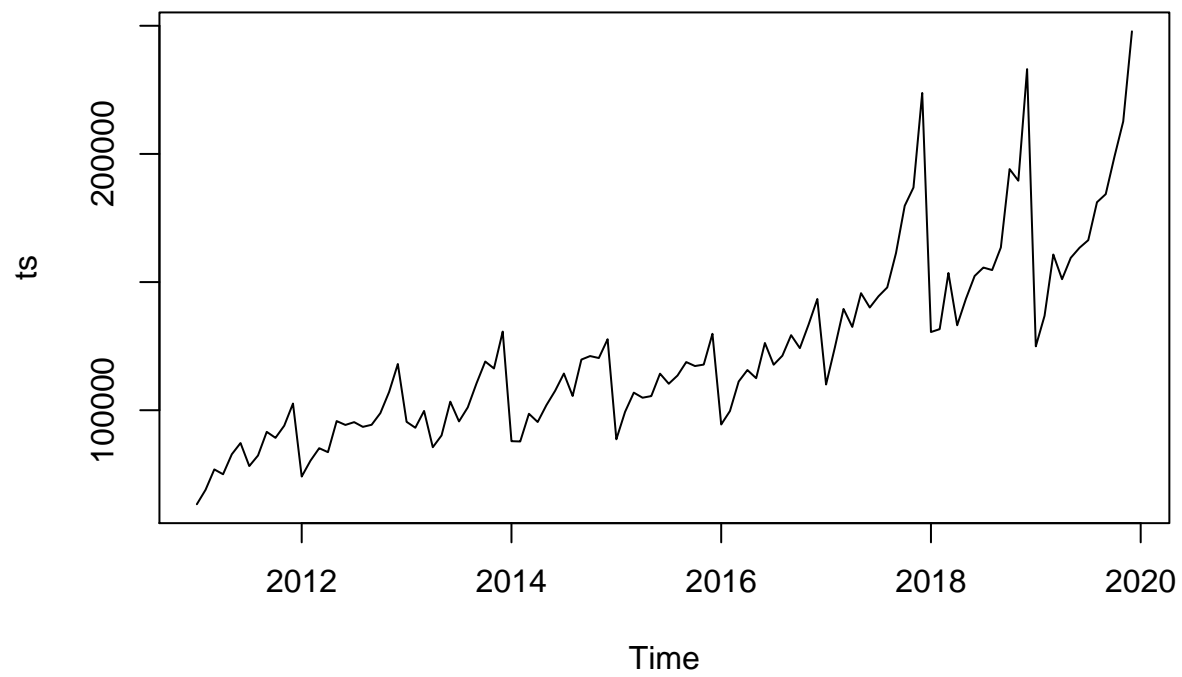
```
plotForecastErrors(ts_forcaste2$residuals)
```



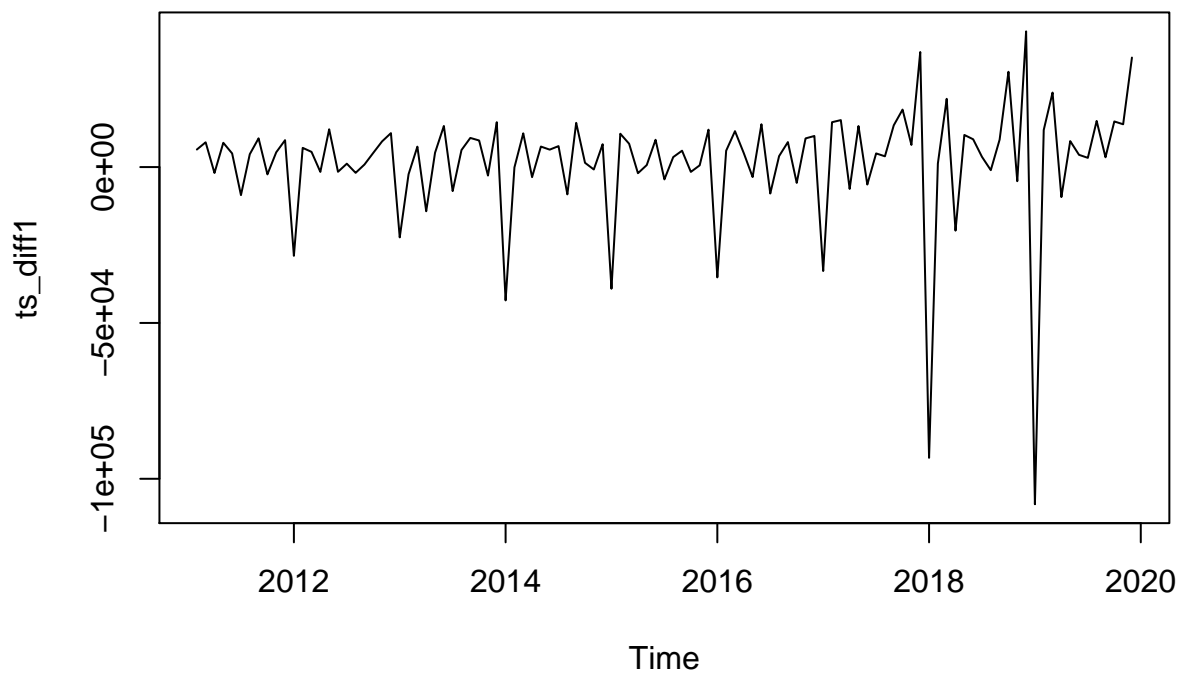
From the time plot, it appears plausible that the forecast errors have constant variance over time. From the histogram of forecast errors, it seems plausible that the forecast errors are normally distributed with mean zero.

Thus, there is little evidence of autocorrelation at lags 1-20 for the forecast errors, and the forecast errors appear to be normally distributed with mean zero and constant variance over time. This suggests that Holt-Winters exponential smoothing provides an adequate predictive model of the log of total productivity, which probably cannot be improved upon. Furthermore, the assumptions upon which the prediction intervals were based are probably valid.

```
plot.ts(ts)
```

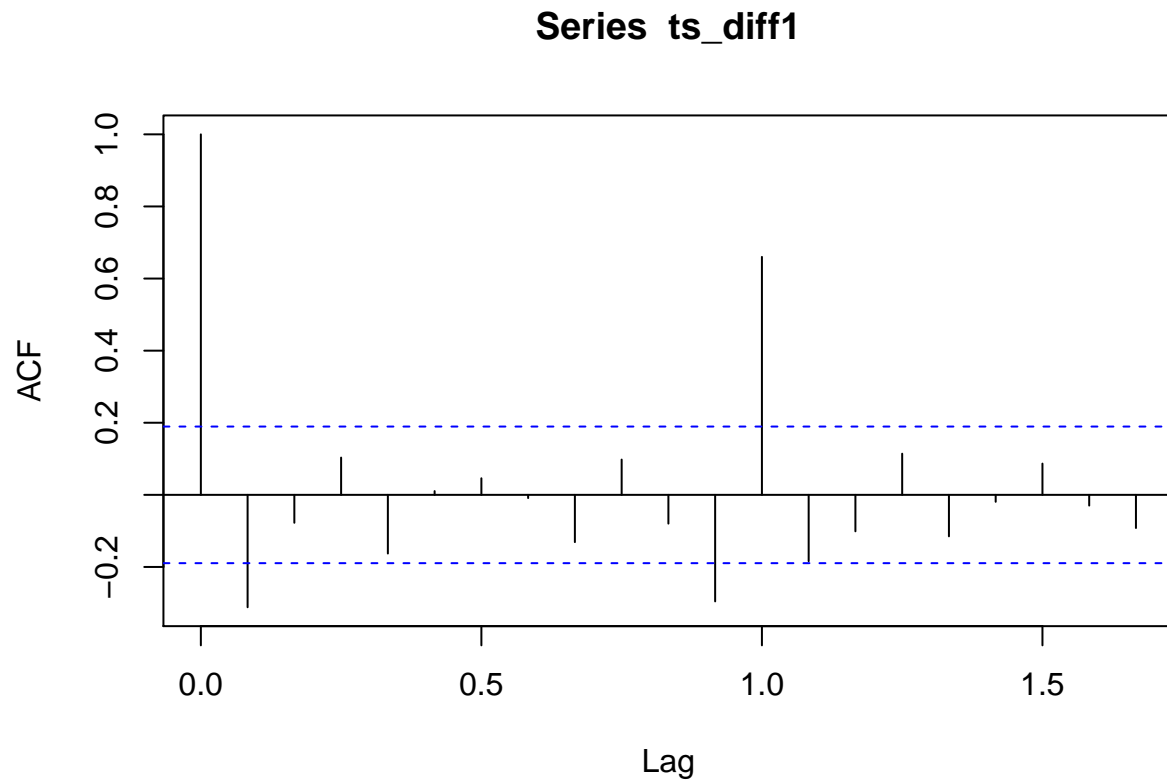


```
ts_diff1 <- diff(ts, differences = 1)
plot.ts(ts_diff1)
```



The time series of differences (above) does appear to be stationary in mean and variance, as the level of the series stays roughly constant over time, and the variance of the series appears roughly constant over time

```
acf(ts_diff1, lag.max=20) # plot a correlogram
```

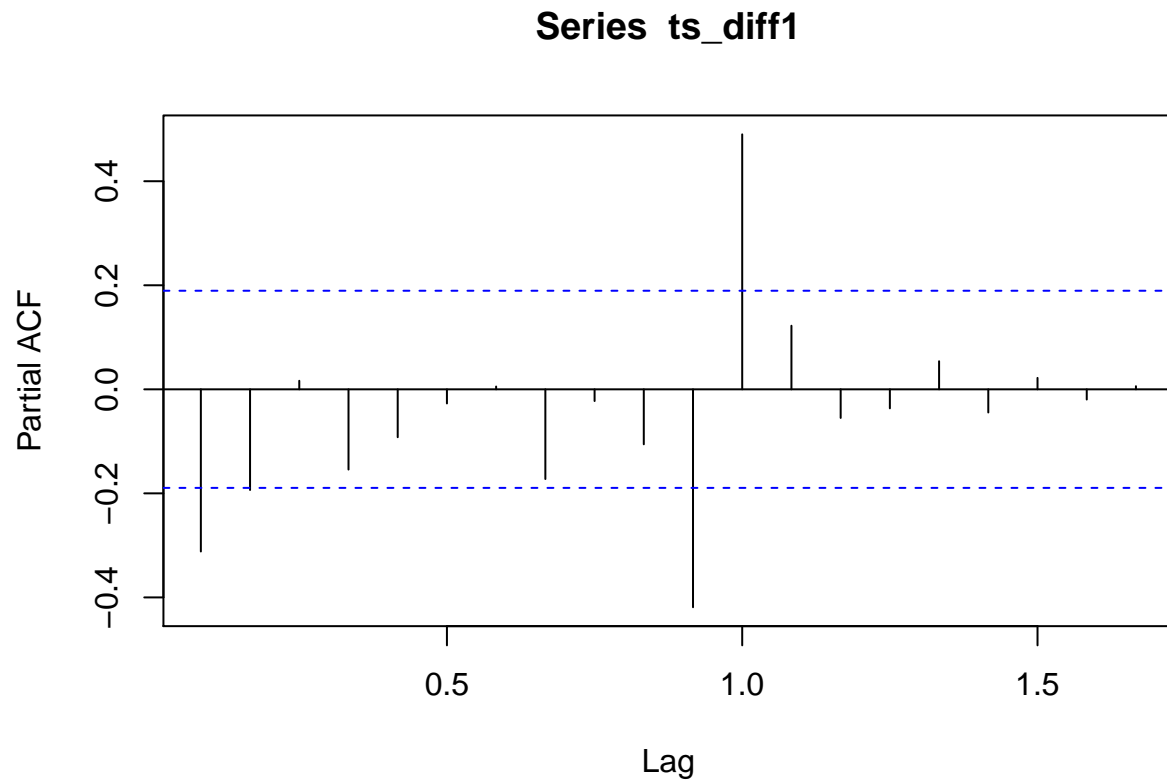


We see from the correlogram that the autocorrelation exceeds the significance bound 3 times but all the others do not exceed

```
acf(ts_diff1, lag.max=20, plot=FALSE) # get the autocorrelation values
```

```
##
## Autocorrelations of series 'ts_diff1', by lag
##
## 0.0000 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333
## 1.000 -0.312 -0.078 0.103 -0.163 0.010 0.046 -0.009 -0.131 0.098 -0.080
## 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
## -0.296 0.660 -0.184 -0.101 0.114 -0.115 -0.019 0.087 -0.030 -0.092
```

```
pacf(ts_diff1, lag.max=20) # plot a partial correlogram
```



```
pacf(ts_diff1, lag.max=20, plot=FALSE) # get the partial autocorrelation values
```

```
##
## Partial autocorrelations of series 'ts_diff1', by lag
##
## 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333 0.9167
## -0.312 -0.193  0.016 -0.154 -0.092 -0.027  0.005 -0.172 -0.023 -0.106 -0.419
## 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
##  0.490  0.122 -0.055 -0.037  0.054 -0.045  0.022 -0.020  0.006
```

Arima, 1,1,1

```
ts_arima = Arima(ts, order=c(1,1,1),seasonal = list(order = c(1,1,1)))
ts_arima
```

```
## Series: ts
## ARIMA(1,1,1)(1,1,1)[12]
##
## Coefficients:
##          ar1          ma1          sar1          sma1
##          0.5303      -0.9098      0.3019      -0.5237
## s.e.      0.1710      0.1144      0.3137      0.2882
```

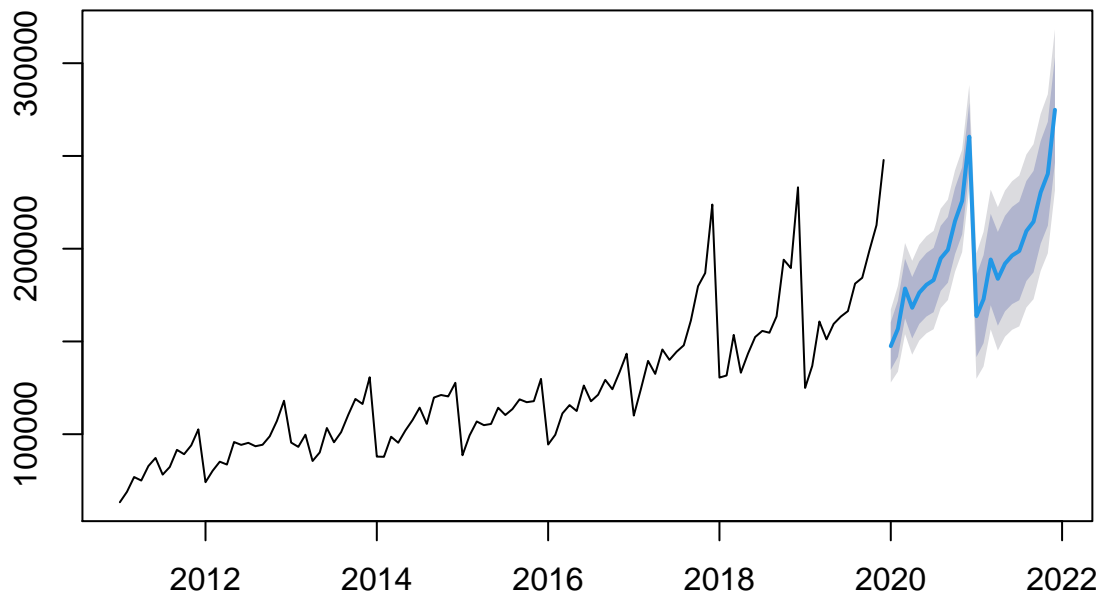
```
##
## sigma^2 estimated as 100477930:  log likelihood=-1008.86
## AIC=2027.73   AICc=2028.4   BIC=2040.5
```

```
ts_arima_forecast = forecast(ts_arima,h = 24)
ts_arima_forecast
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2020	147503.2	134657.0	160349.3	127856.7	167149.7
## Feb 2020	156812.5	141694.0	171930.9	133690.8	179934.1
## Mar 2020	178520.7	162471.5	194570.0	153975.5	203066.0
## Apr 2020	168153.4	151609.4	184697.3	142851.6	193455.2
## May 2020	176316.3	159448.7	193183.9	150519.6	202113.0
## Jun 2020	180548.0	163432.5	197663.4	154372.1	206723.8
## Jul 2020	183095.0	165768.1	200421.8	156595.8	209594.1
## Aug 2020	194784.4	177265.1	212303.8	167990.9	221578.0
## Sep 2020	199379.3	181677.9	217080.7	172307.4	226451.3
## Oct 2020	214902.8	197025.5	232780.1	187561.8	242243.8
## Nov 2020	225786.7	207737.4	243836.0	198182.7	253390.7
## Dec 2020	260407.7	242189.2	278626.2	232544.9	288270.4
## Jan 2021	163636.7	141560.0	185713.4	129873.4	197400.1
## Feb 2021	172790.7	149071.5	196509.9	136515.3	209066.0
## Mar 2021	194167.2	169530.8	218803.7	156489.0	231845.5
## Apr 2021	183753.9	158490.7	209017.2	145117.1	222390.8
## May 2021	191957.1	166198.7	217715.6	152563.0	231351.3
## Jun 2021	196330.2	170141.8	222518.7	156278.5	236382.0
## Jul 2021	198777.1	172193.5	225360.8	158120.9	239433.4
## Aug 2021	209540.0	182581.0	236498.9	168309.8	250770.2
## Sep 2021	214579.2	187257.4	241901.0	172794.1	256364.2
## Oct 2021	230371.7	202695.6	258047.7	188044.8	272698.5
## Nov 2021	240397.6	212373.7	268421.4	197538.8	283256.4
## Dec 2021	274863.2	246496.8	303229.6	231480.6	318245.8

```
forecast::plot.forecast(ts_arima_forecast)
```

Forecasts from ARIMA(1,1,1)(1,1,1)[12]



Growth

```
# this_year_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1]

# year_2019_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1][c(1:2),]
# sum_year_2019 = sum(c(year_2019,year_2019_predict_ARIMA))
# year_2020 = (as.data.frame(ts_arma_forecast))[1][c(3:14),]
year_2020 <- (as.data.frame(ts_arma_forecast))[1][c(1:12),]
year_2021 <- (as.data.frame(ts_arma_forecast))[1][c(13:24),]

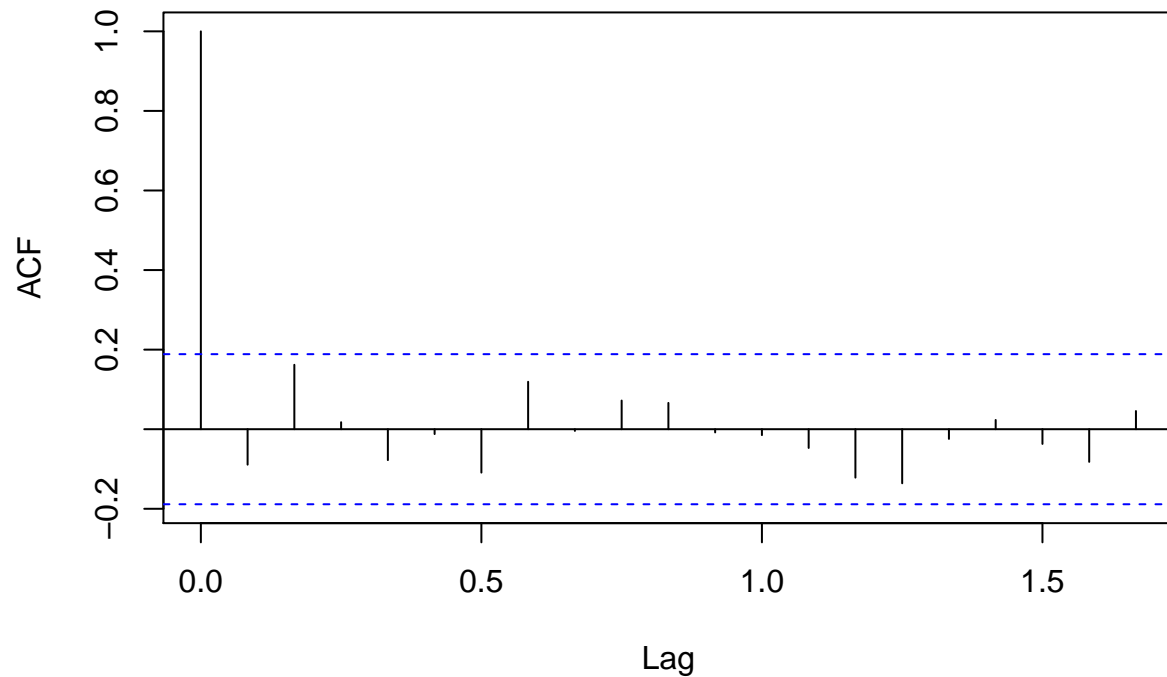
growth_ARIMA_21 <- growth(sum(year_2021), sum(year_2020))
growth_ARIMA_20 <- growth(sum(year_2020), sum(year_2019))
```

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arma_forecast$residuals, lag.max=20)
```


Series ts_arma_forecast\$residuals

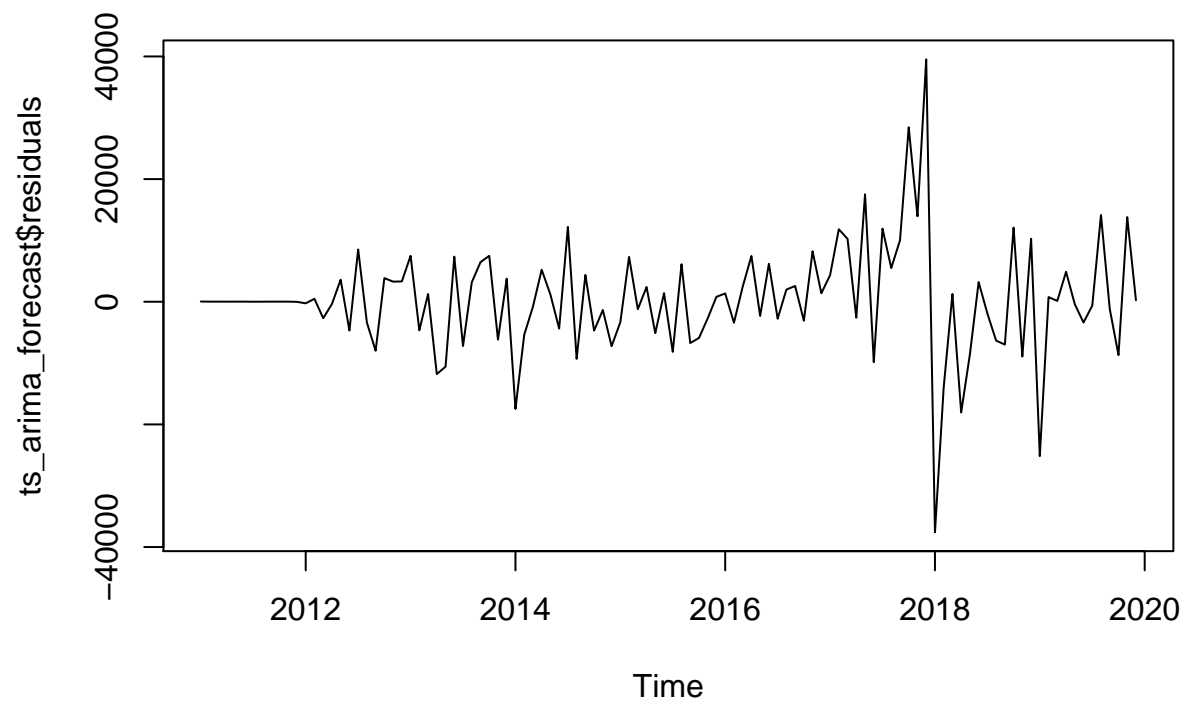


```
Box.test(ts_arma_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: ts_arma_forecast$residuals  
## X-squared = 14.835, df = 20, p-value = 0.7858
```

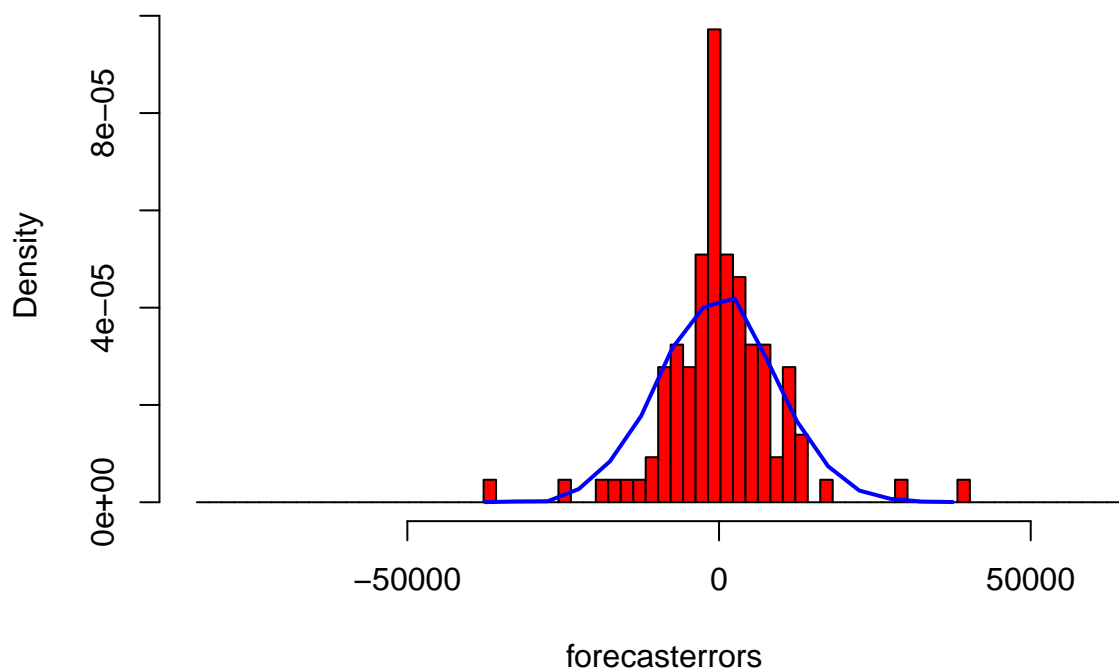
we can reject the null hypothesis, it's rather similar to the HW

```
plot.ts(ts_arma_forecast$residuals)           # make time plot of forecast errors
```



```
plotForecastErrors(ts_arma_forecast$residuals)
```

Histogram of forecasterrors



```
# Arima, 0,1,0 as given from the loop
```

```
ts_arima = Arima(ts, order=c(2,1,1),seasonal = list(order = c(2,1,0)))
ts_arima
```

```
## Series: ts
## ARIMA(2,1,1)(2,1,0)[12]
##
## Coefficients:
##          ar1      ar2      ma1      sar1      sar2
##          0.5014  0.1847 -0.9647 -0.1801 -0.1210
## s.e.      0.1104  0.1095   0.0545   0.1072   0.1275
##
## sigma^2 estimated as 98212454: log likelihood=-1007.38
## AIC=2026.76   AICc=2027.71   BIC=2042.08
```

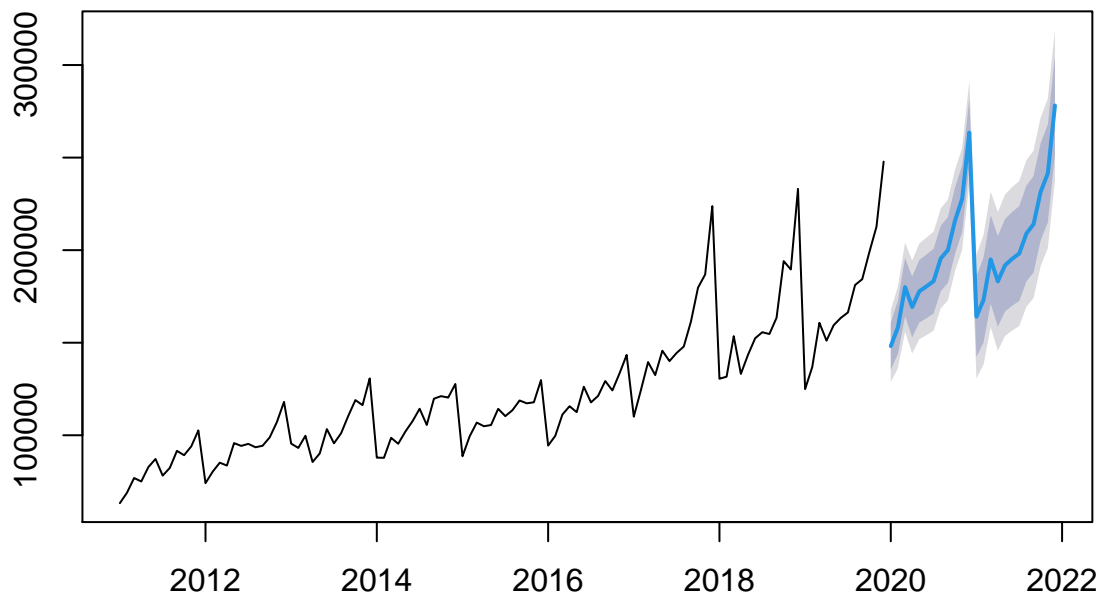
```
ts_arima_forecast = forecast(ts_arima,h = 24)
ts_arima_forecast
```

```
##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## Jan 2020      148158.0 135457.0 160858.9 128733.5 167582.4
## Feb 2020      158278.8 143863.5 172694.1 136232.5 180325.1
## Mar 2020      180009.1 164311.4 195706.7 156001.6 204016.6
## Apr 2020      169295.4 152873.1 185717.8 144179.6 194411.3
## May 2020      177769.1 160862.8 194675.5 151913.1 203625.2
## Jun 2020      180412.5 163176.7 197648.4 154052.5 206772.6
```

```
## Jul 2020      183251.2 165777.6 200724.9 156527.6 209974.9
## Aug 2020      195512.8 177859.8 213165.7 168514.9 222510.6
## Sep 2020      200057.6 182263.4 217851.8 172843.7 227271.5
## Oct 2020      215970.9 198060.8 233881.1 188579.7 243362.1
## Nov 2020      227724.5 209715.6 245733.3 200182.4 255266.6
## Dec 2020      263501.5 245406.0 281596.9 235826.9 291176.1
## Jan 2021      164047.6 142279.7 185815.4 130756.4 197338.7
## Feb 2021      173142.9 150211.5 196074.4 138072.3 208213.5
## Mar 2021      194991.7 171110.3 218873.2 158468.3 231515.2
## Apr 2021      183147.4 158642.3 207652.5 145670.1 220624.7
## May 2021      191818.4 166846.9 216789.9 153627.8 230009.0
## Jun 2021      195284.0 169955.0 220613.0 156546.7 234021.4
## Jul 2021      198175.2 172558.1 223792.2 158997.3 237353.0
## Aug 2021      208976.0 183118.1 234833.9 169429.8 248522.2
## Sep 2021      213949.9 187883.6 240016.3 174084.8 253815.0
## Oct 2021      231559.6 205307.5 257811.7 191410.5 271708.7
## Nov 2021      241456.2 215034.7 267877.8 201048.0 281864.5
## Dec 2021      278133.8 251554.5 304713.0 237484.4 318783.2
```

```
forecast::plot.forecast(ts_arima_forecast)
```

Forecasts from ARIMA(2,1,1)(2,1,0)[12]



```
## Growth
```

```
# this_year_predict_ARIMA <- (as.data.frame(ts_arima_forecast))[1]
#
# year_2019_predict_ARIMA <- (as.data.frame(ts_arima_forecast))[1][c(1:2),]
```

```
# sum_year_2019 = sum(c(year_2019,year_2019_predict_ARIMA))
# year_2020 = (as.data.frame(ts_arima_forecast))[1][c(3:14),]

year_2020 <- (as.data.frame(ts_arima_forecast))[1][c(1:12),]
year_2021 <- (as.data.frame(ts_arima_forecast))[1][c(13:24),]

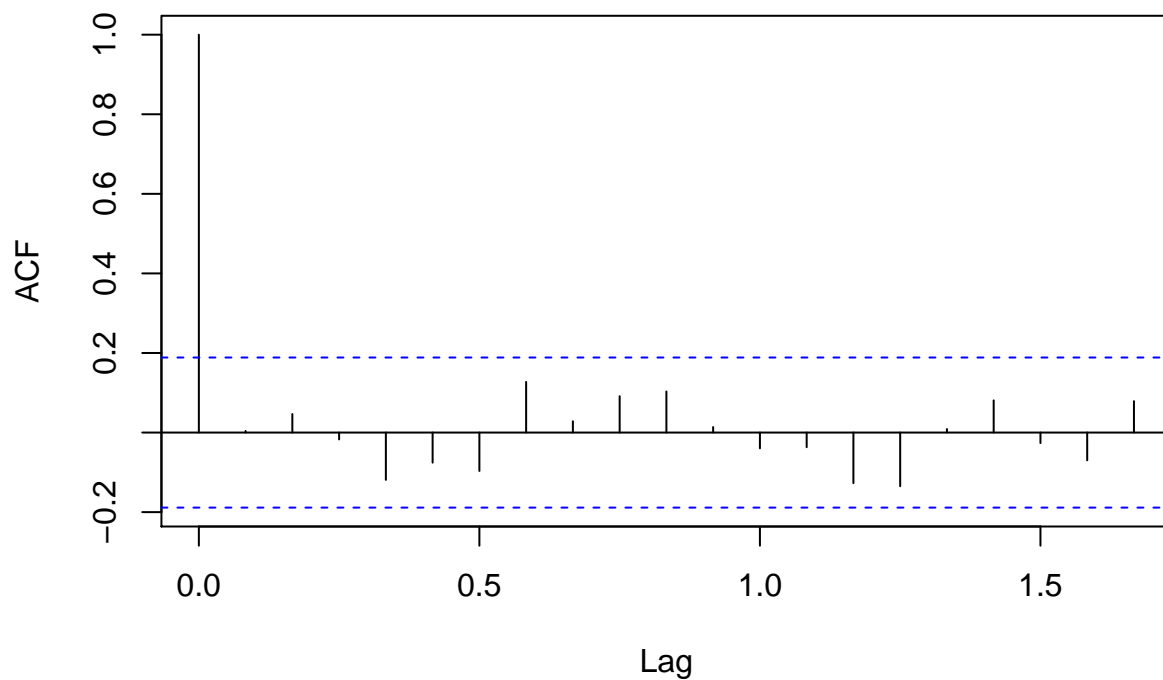
growth_ARIMA2_21 <- growth(sum(year_2021), sum(year_2020))
growth_ARIMA2_20 <- growth(sum(year_2020), sum(year_2019))
```

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arima_forecast$residuals, lag.max=20)
```

Series ts_arima_forecast\$residuals

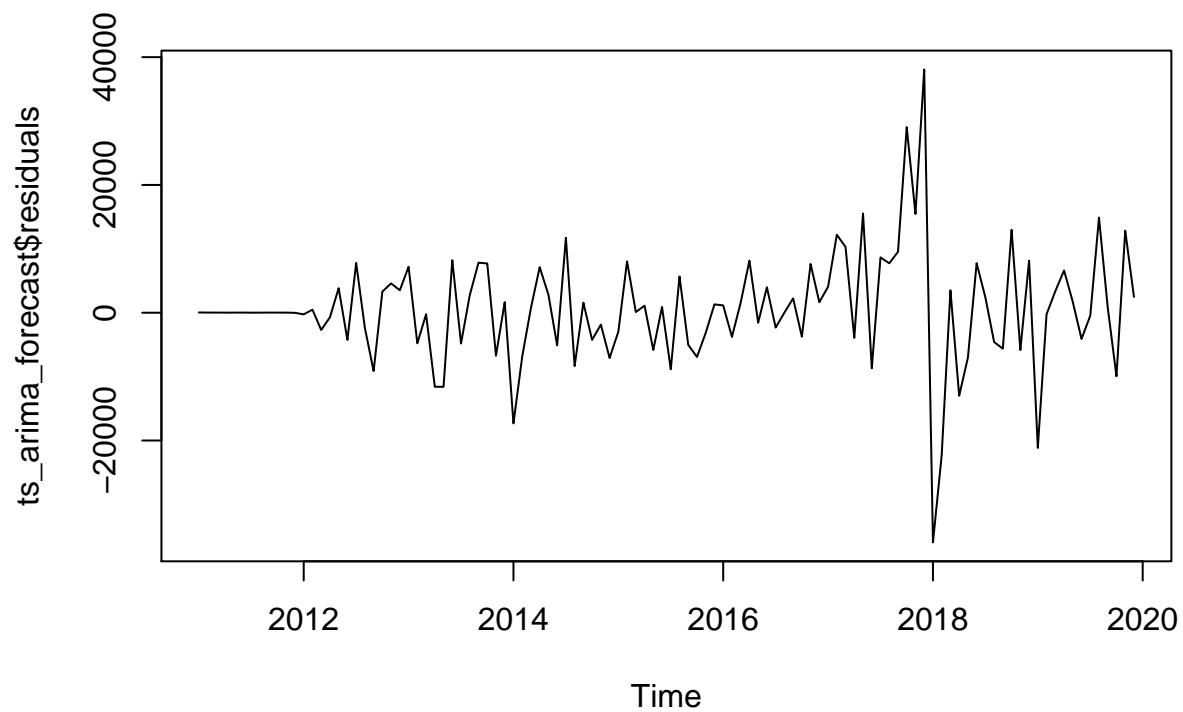


```
Box.test(ts_arima_forecast$residuals, lag=20, type="Ljung-Box")
```

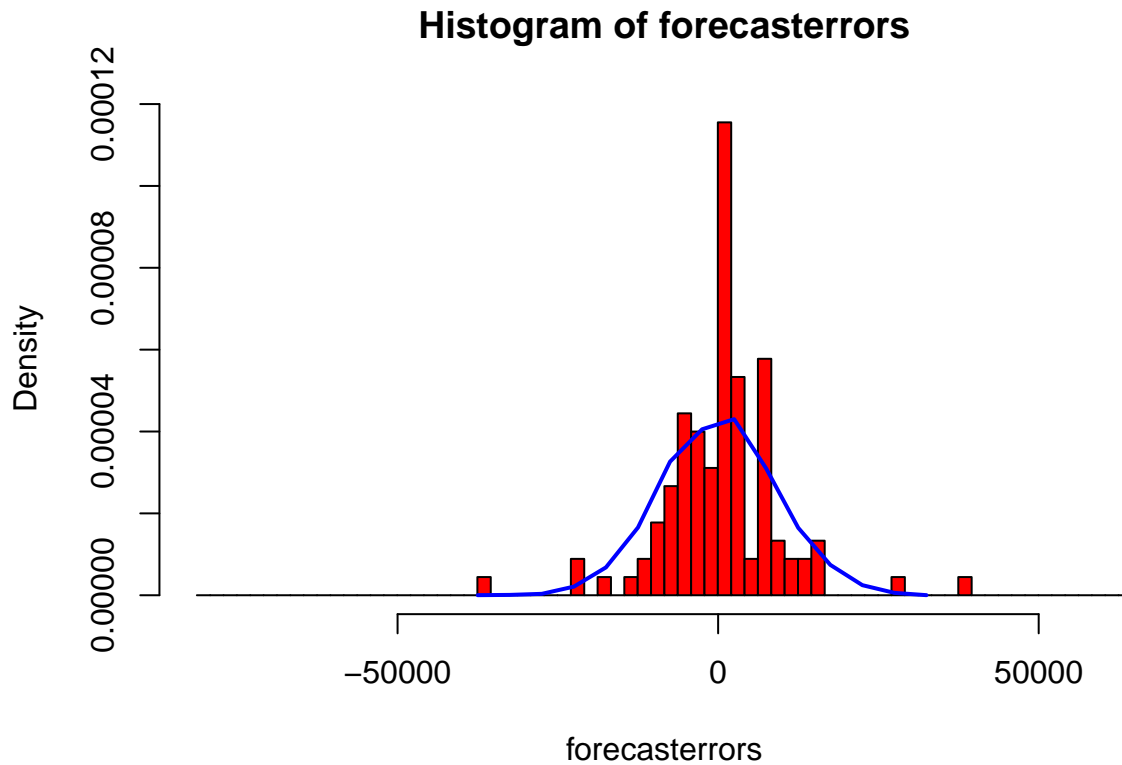
```
##
## Box-Ljung test
##
## data: ts_arima_forecast$residuals
## X-squared = 15.147, df = 20, p-value = 0.768
```

we can reject the null hypothesis, it's rather similar to the HW

```
plot.ts(ts_arima_forecast$residuals)           # make time plot of forecast errors
```



```
plotForecastErrors(ts_arima_forecast$residuals)
```



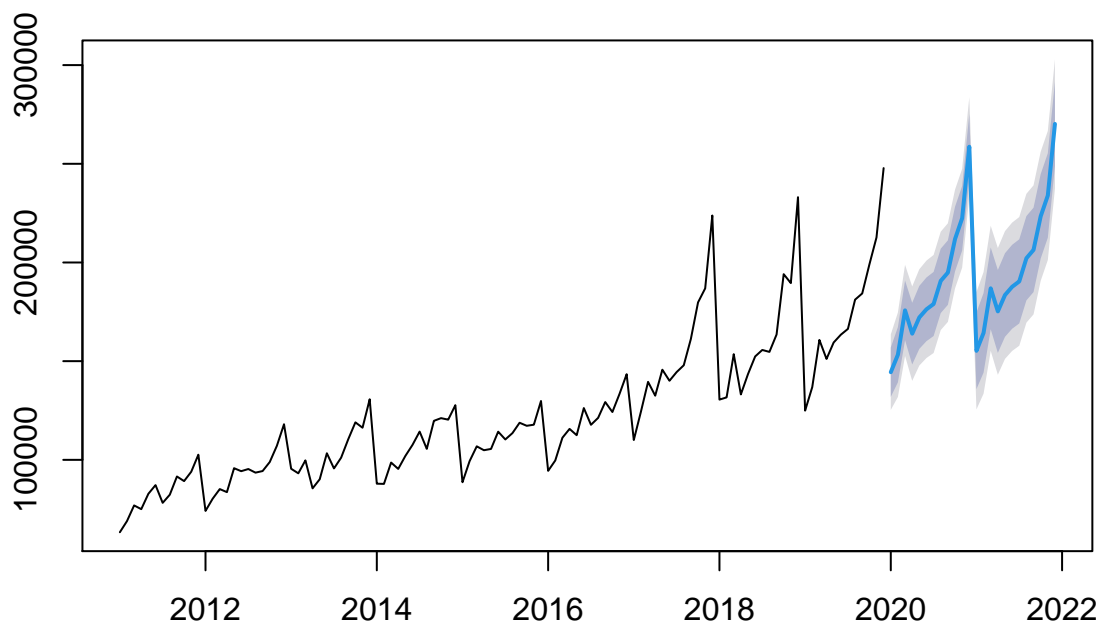
A model chosen automatically

```
fit <- auto.arima(ts,max.p = 5,max.q = 5,max.P = 5,max.Q = 5,max.d = 3,seasonal = TRUE)
fit
```

```
## Series: ts
## ARIMA(2,0,0)(0,1,1)[12] with drift
##
## Coefficients:
##          ar1      ar2      sma1      drift
##          0.5027  0.1991  -0.1932  966.1516
## s.e.  0.0991  0.0995   0.1129  219.5238
##
## sigma^2 estimated as 95095303:  log likelihood=-1016.47
## AIC=2042.93   AICc=2043.6   BIC=2055.75
```

```
fit_forecast = forecast(fit,h=24)
plot(fit_forecast)
```

Forecasts from ARIMA(2,0,0)(0,1,1)[12] with drift



```
# str(fit)
```

Growth

```
# year_2021_predict_auto.arima <- (as.data.frame(fit_forecast))[1]
# year_2021_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4]
# year_2021_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5]
#
# growth_auto.arima <- growth(sum(year_2021_predict_auto.arima), sum(year_2020))
# growth_auto.arima_95_low <- growth(sum(year_2021_predict_auto.arima_95_low), sum(year_2020))
# growth_auto.arima_95_high <- growth(sum(year_2021_predict_auto.arima_95_high), sum(year_2020))
#
# growth_auto.arima
# growth_auto.arima_95_low
# growth_auto.arima_95_high

year_2020 <- (as.data.frame(fit_forecast))[1][c(1:12),]
year_2021 <- (as.data.frame(fit_forecast))[1][c(13:24),]

growth_auto.arima_21 <- growth(sum(year_2021), sum(year_2020))
growth_auto.arima_20 <- growth(sum(year_2020), sum(year_2019))
```


all the growths

```
# growth_ARIMA = -growth_ARIMA
#
# growth_ARIMA2 = -growth_ARIMA2
#
# growth_auto.arima
#
# growth_HW

growth_ARIMA_20
```

```
## [1] 0.09503449
```

```
growth_ARIMA_21
```

```
## [1] 0.08090016
```

```
growth_ARIMA2_20
```

```
## [1] 0.1016115
```

```
growth_ARIMA2_21
```

```
## [1] 0.07597644
```

```
growth_auto.arima_20
```

```
## [1] 0.07442554
```

```
growth_auto.arima_21
```

```
## [1] 0.06070137
```