

# tsf\_export

*Kevork Sulahian*

*October 28, 2019*

```
library(readxl)
library(forecast)
```

```
# library(readxl)
```

```
df <- read_xlsx("Export_for_TS.xlsx")
```

```
## New names:
## * `` -> ...2
```

```
df = df[1,]
df = df[-c(1,2)]
```

```
df2 = read_xlsx('export_19xlsx.xlsx')
```


```
## New names:
## * `` -> ...1
```

```
df = t(df)
df2$...1 = c(paste0(2019,"-",1:11))
rownames(df2) =df2$...1
```

```
## Warning: Setting row names on a tibble is deprecated.
```


```
df2$...1 = NULL
colnames(df2) = "V1"
df = as.data.frame(df)
```

```
df3 = rbind(df,df2)
ts = ts(df3,start=c(2009,1), frequency = c(12))
```



In this case, it appears that an additive model is not appropriate for describing this time series, since the size of the seasonal fluctuations and random fluctuations seem to increase with the level of the time series. Thus, we may need to transform the time series in order to get a transformed time series that can be described using an additive model. For example, we can transform the time series by calculating the natural log of the original data:

```
log_ts <- log(ts)
plot.ts(log_ts)
```



## Decomposing Time Series

Decomposing a time series means separating it into its constituent components, which are usually a trend component and an irregular component, and if it is a seasonal time series, a seasonal component.

### Decomposing Seasonal Data

A seasonal time series consists of a trend component, a seasonal component and an irregular component. Decomposing the time series means separating the time series into these three components: that is, estimating these three components.

```
ts_components <- decompose(ts)
```


we can print out the estimated values of the seasonal component

```
ts_components$seasonal
```

##	Jan	Feb	Mar	Apr	May	Jun
## 2009	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
## 2010	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
## 2011	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
## 2012	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
## 2013	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528

##	2014	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
##	2015	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
##	2016	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
##	2017	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
##	2018	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
##	2019	-35872.943	-5113.997	579.187	-4874.741	-7357.388	8874.528
##		Jul	Aug	Sep	Oct	Nov	Dec
##	2009	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2010	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2011	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2012	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2013	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2014	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2015	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2016	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2017	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2018	3440.638	2149.841	8722.117	8915.671	8744.667	11792.421
##	2019	3440.638	2149.841	8722.117	8915.671	8744.667	


## Decomposition of additive time series



The plot above shows the original time series (top), the estimated trend component (second from top), the estimated seasonal component (third from top), and the estimated irregular component (bottom)

## Seasonally Adjusting

```
ts_seasonall <- ts - ts_components$seasonal
```



## Holt-Winters Exponential Smoothing

```
ts_forcaste <- HoltWinters(ts)
ts_forcaste
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = ts)
##
## Smoothing parameters:
##  alpha: 0.3631002
##  beta : 0.00776954
##  gamma: 0.3604275
##
## Coefficients:
##           [,1]
## a    241120.5637
## b      1841.9013
## s1     6407.2267
```

```
## s2 -45300.7926
## s3 -7832.0311
## s4 -255.7858
## s5 -8230.8329
## s6 -12108.1995
## s7 5114.0000
## s8 6623.6290
## s9 5748.0034
## s10 4318.7619
## s11 7502.7160
## s12 7565.4575
```


```
#
```

The value of alpha (0.35) is relatively low, indicating that the estimate of the level at the current time point is based upon both recent observations and some observations in the more distant past. The value of beta is 0.01, indicating that the estimate of the slope  $b$  of the trend component is updated but doesn't have much effect over the time series, and instead is set equal to its initial value. This makes good intuitive sense, as the level changes quite a bit over the time series, but the slope  $b$  of the trend component remains roughly the same. In contrast, the value of gamma (0.38) is high, indicating that the estimate of the seasonal component at the current time point is not just based upon very recent observations

```
ts_forecaste$SSE
```

```
## [1] 22399824125
```


## Holt-Winters filtering



```
ts_forcaste2 = forecast::forecast.HoltWinters(ts_forcaste, h= 13)
(as.data.frame(ts_forcaste2))[1]
```

```
##          Point Forecast
## Dec 2019      249369.7
## Jan 2020      199503.6
## Feb 2020      238814.2
## Mar 2020      248232.4
## Apr 2020      242099.2
## May 2020      240063.8
## Jun 2020      259127.9
## Jul 2020      262479.4
## Aug 2020      263445.7
## Sep 2020      263858.3
## Oct 2020      268884.2
## Nov 2020      270788.8
## Dec 2020      271472.5
```

## Forecasts from HoltWinters




```
## Growth
```

```
year_2019 <- window(ts, 2019)
year_2019_predict_HW <- (as.data.frame(ts_forcaste2))[1][c(1),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_HW))
year_2020 = (as.data.frame(ts_forcaste2))[1][c(2:13),]
growth_HW <- growth(sum(year_2020),sum_year_2019)
growth_HW
```

```
## [1] 0.1508492
```

We can investigate whether the predictive model can be improved upon by checking whether the in-sample forecast errors show non-zero autocorrelations at lags 1-20, by making a correlogram and carrying out the Ljung-Box test:




```
##  
## Box-Ljung test  
##  
## data: ts_forcaste2$residuals  
## X-squared = 13.33, df = 20, p-value = 0.8628
```

The correlogram shows that the autocorrelations for the in-sample forecast errors do not exceed the significance bounds for lags 1-20. Furthermore, the p-value for Ljung-Box test is 0.9, indicating that there is no evidence of non-zero autocorrelations at lags 1-20.

We can check whether the forecast errors have constant variance over time, and are normally distributed with mean zero, by making a time plot of the forecast errors and a histogram (with overlaid normal curve):

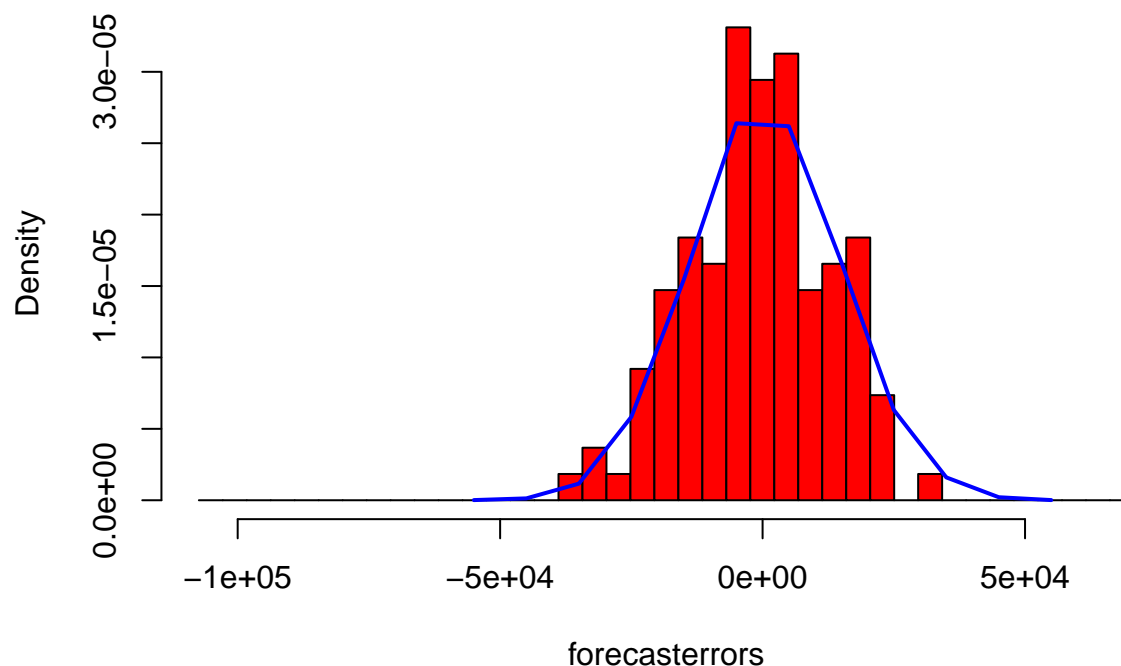
```
plot.ts(ts_forcaste2$residuals)
```





```
plotForecastErrors(ts_forcaste2$residuals)
```


## Histogram of forecasterrors




From the time plot, it appears plausible that the forecast errors have constant variance over time. From the histogram of forecast errors, it seems plausible that the forecast errors are normally distributed with mean zero.

Thus, there is little evidence of autocorrelation at lags 1-20 for the forecast errors, and the forecast errors appear to be normally distributed with mean zero and constant variance over time. This suggests that Holt-Winters exponential smoothing provides an adequate predictive model of the log of total productivity, which probably cannot be improved upon. Furthermore, the assumptions upon which the prediction intervals were based are probably valid.

```
plot.ts(ts)
```




```
ts_diff1 <- diff(ts, differences = 1)
plot.ts(ts_diff1)
```



The time series of differences (above) does appear to be stationary in mean and variance, as the level of the series stays roughly constant over time, and the variance of the series appears roughly constant over time

```
acf(ts_diff1, lag.max=20) # plot a correlogram
```




We see from the correlogram that the autocorrelation exceeds the significance bound 3 times but all the others do not exceed

```
acf(ts_diff1, lag.max=20, plot=FALSE) # get the autocorrelation values
```

```
##
## Autocorrelations of series 'ts_diff1', by lag
##
## 0.0000 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500
## 1.000 -0.442 -0.055 0.035 0.056 -0.083 0.097 -0.200 0.122 0.092
## 0.8333 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833
## -0.132 -0.204 0.484 -0.165 -0.168 0.149 -0.019 -0.034 0.030 -0.119
## 1.6667
## 0.033
```

```
pacf(ts_diff1, lag.max=20) # plot a partial correlogram
```



```
pacf(ts_diff1, lag.max=20, plot=FALSE) # get the partial autocorrelation values
```

```
##
## Partial autocorrelations of series 'ts_diff1', by lag
##
## 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333
## -0.442 -0.311 -0.185 -0.038 -0.085 0.050 -0.202 -0.095 0.077 -0.039
## 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
## -0.354 0.241 0.232 -0.126 0.051 0.095 -0.035 -0.121 0.022 -0.068
```

**Arima, 0,1,0**


```
ts_arma = Arima(ts, order=c(0,1,0),seasonal = list(order = c(0,1,0)))
ts_arma
```

```
## Series: ts
## ARIMA(0,1,0)(0,1,0)[12]
##
## sigma^2 estimated as 442379690: log likelihood=-1341.99
## AIC=2685.98 AICc=2686.01 BIC=2688.75
```

```
ts_arma_forecast = forecast(ts_arma,h = 13)
ts_arma_forecast
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Dec 2019	263021.4	236066.8	289976.1	221797.9	304245.0
## Jan 2020	200410.5	162290.9	238530.2	142111.6	258709.5
## Feb 2020	247610.5	200923.7	294297.4	176209.2	319011.9
## Mar 2020	261510.5	207601.2	315419.8	179063.4	343957.7
## Apr 2020	262910.5	202638.1	323183.0	170731.8	355089.3
## May 2020	263310.5	197285.4	329335.7	162333.8	364287.3
## Jun 2020	267310.5	195995.2	338625.8	158243.2	376377.9
## Jul 2020	290510.5	214271.3	366749.8	173912.6	407108.4
## Aug 2020	294610.5	213746.6	375474.5	170939.8	418281.3
## Sep 2020	297810.5	212572.4	383048.6	167450.1	428170.9
## Oct 2020	288810.5	199412.1	378209.0	152087.4	425533.7
## Nov 2020	312710.5	219336.9	406084.2	169907.9	455513.2
## Dec 2020	318032.0	210213.4	425850.6	153137.7	482926.3

```
forecast::plot.forecast(ts_arma_forecast)
```



Growth

```

this_year_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1]

# growth_ARIMA <- growth(sum(c(this_year,as.numeric(this_year_predict_ARIMA$`Point Forecast`))), sum(la
# growth_ARIMA

year_2019_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1][c(1),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_ARIMA))
year_2020 = (as.data.frame(ts_arma_forecast))[1][c(2:13),]
growth_ARIMA <- growth(sum_year_2019, sum(year_2020))
-growth_ARIMA

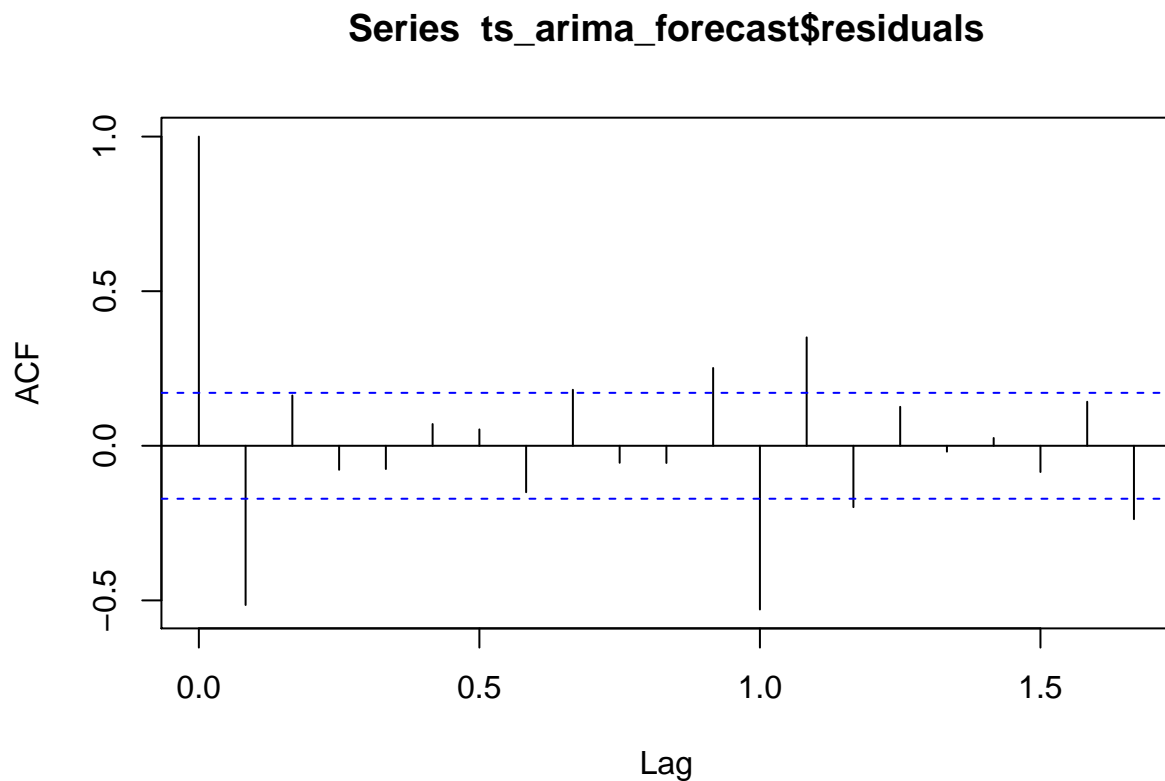
```

```
## [1] 0.1997026
```

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arma_forecast$residuals, lag.max=20)
```






```
Box.test(ts_arima_forecast$residuals, lag=20, type="Ljung-Box")
```


```
##  
## Box-Ljung test  
##  
## data: ts_arima_forecast$residuals  
## X-squared = 140.34, df = 20, p-value < 2.2e-16
```

p value too high to reject

```
plot.ts(ts_arima_forecast$residuals) # make time plot of forecast errors
```



```
plotForecastErrors(ts_arima_forecast$residuals)
```




## A model chosen automatically

```
fit3 <- auto.arima(ts)
fit3
```

```
## Series: ts
## ARIMA(1,0,1)(0,1,1)[12] with drift
##
## Coefficients:
##          ar1          ma1          sma1          drift
##          0.9554   -0.5899   -0.8838   1299.0445
## s.e.    0.0387    0.0854    0.1439    204.5971
##
## sigma^2 estimated as 160329918:  log likelihood=-1299.54
## AIC=2609.07   AICc=2609.6   BIC=2622.97
```

```
fit_forecast = forecast(fit3,h=15)
plot(fit_forecast)
```

## Forecasts from ARIMA(1,0,1)(0,1,1)[12] with drift



```
# str(fit)
```

## Growth

```
year_2019 <- window(ts, 2019)
year_2019_predict_HW <- (as.data.frame(ts_forcaste2))[1][c(1),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_HW))
year_2020 = (as.data.frame(ts_forcaste2))[1][c(2:13),]

year_2019_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(1),]
year_2019_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(1),]
year_2019_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(1),]

sum_year_2019 = sum(c(year_2019,year_2019_predict_auto.arima))
sum_year_2019_low = sum(c(year_2019,year_2019_predict_auto.arima_95_low))
sum_year_2019_high = sum(c(year_2019,year_2019_predict_auto.arima_95_high))

year_2020_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(2:13),]
year_2020_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(2:13),]
year_2020_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(2:13),]
```

```
growth_auto.arima <- growth(sum(year_2020_predict_auto.arima),sum_year_2019)
growth_auto.arima_95_low <- growth(sum(year_2020_predict_auto.arima_95_low),sum_year_2019_low)
growth_auto.arima_95_high <- growth(sum(year_2020_predict_auto.arima_95_high),sum_year_2019_high)

growth_auto.arima
```

```
## [1] 0.1056297
```

```
growth_auto.arima_95_low
```

```
## [1] -0.0307078
```

```
growth_auto.arima_95_high
```

```
## [1] 0.2393957
```