

# tsf\_export

Kevork Sulahian

October 28, 2019

```
library(readxl)
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```

```
# library(readxl)
```

```
df <- read_xlsx("Export_for_TS.xlsx")
```

```
## New names:
## * ' ' -> ...2
```

```
df = df[1,]
df = df[-c(1,2)]
```

```
df2 = read_xlsx('export_19xlsx.xlsx')
```


```
## New names:
## * ' ' -> ...1
```

```
df = t(df)
# df2$...1 = c(paste0(2019,"-",1:12))
rownames(df2) =df2$...1
```

```
## Warning: Setting row names on a tibble is deprecated.
```


```
df2$...1 = NULL
colnames(df2) = "V1"
df = as.data.frame(df)
```

```
df3 = rbind(df,df2)
ts = ts(df3,start=c(2009,1), frequency = c(12))
```



In this case, it appears that an additive model is not appropriate for describing this time series, since the size of the seasonal fluctuations and random fluctuations seem to increase with the level of the time series. Thus, we may need to transform the time series in order to get a transformed time series that can be described using an additive model. For example, we can transform the time series by calculating the natural log of the original data:

```
log_ts <- log(ts)
plot.ts(log_ts)
```



## ##Decomposing Time Series

Decomposing a time series means separating it into its constituent components, which are usually a trend component and an irregular component, and if it is a seasonal time series, a seasonal component.

###Decomposing Seasonal Data A seasonal time series consists of a trend component, a seasonal component and an irregular component. Decomposing the time series means separating the time series into these three components: that is, estimating these three components.

```
ts_components <- decompose(ts)
```


we can print out the estimated values of the seasonal component

```
ts_components$seasonal
```

##	Jan	Feb	Mar	Apr	May	Jun
## 2009	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2010	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2011	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2012	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2013	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2014	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2015	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2016	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2017	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602
## 2018	-35868.9712	-5110.0251	583.1592	-4870.7687	-7353.4162	7568.9602

```
## 2019 -35868.9712 -5110.0251 583.1592 -4870.7687 -7353.4162 7568.9602
## 2020 -35868.9712
##      Jul      Aug      Sep      Oct      Nov      Dec
## 2009  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2010  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2011  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2012  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2013  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2014  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2015  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2016  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2017  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2018  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2019  4706.4842  2153.8132  8726.0894  8919.6430  8748.6388 11796.3933
## 2020
```


## Decomposition of additive time series



The plot above shows the original time series (top), the estimated trend component (second from top), the estimated seasonal component (third from top), and the estimated irregular component (bottom)

## Seasonally Adjusting

```
ts_seasonall <- ts - ts_components$seasonal
```



## Holt-Winters Exponential Smoothing

```
ts_forcaste <- HoltWinters(ts)
ts_forcaste
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = ts)
##
## Smoothing parameters:
##  alpha: 0.3550801
##  beta : 0.01223594
##  gamma: 0.3718725
##
## Coefficients:
##           [,1]
## a    227306.8042
## b     1573.4174
## s1   -7376.7400
## s2     343.7398
## s3   -7858.7490
## s4  -11901.5253
```

```
## s5      5357.7476
## s6      7156.3868
## s7      6538.0354
## s8      4815.7484
## s9      8180.6545
## s10     8364.9546
## s11     1449.0512
## s12    -50544.5232
```


```
#
```

The value of alpha (0.35) is relatively low, indicating that the estimate of the level at the current time point is based upon both recent observations and some observations in the more distant past. The value of beta is 0.01, indicating that the estimate of the slope  $b$  of the trend component is updated but doesn't have much effect over the time series, and instead is set equal to its initial value. This makes good intuitive sense, as the level changes quite a bit over the time series, but the slope  $b$  of the trend component remains roughly the same. In contrast, the value of gamma (0.38) is high, indicating that the estimate of the seasonal component at the current time point is not just based upon very recent observations

```
ts_forecaste$SSE
```

```
## [1] 23479926065
```


## Forecasts from HoltWinters



```
ts_forcaste2 = forecast::forecast.HoltWinters(ts_forcaste, h= 11)
(as.data.frame(ts_forcaste2))[1]
```

```
##          Point Forecast
## Feb 2020      221503.5
## Mar 2020      230797.4
## Apr 2020      224168.3
## May 2020      221698.9
## Jun 2020      240531.6
## Jul 2020      243903.7
## Aug 2020      244858.8
## Sep 2020      244709.9
## Oct 2020      249648.2
## Nov 2020      251405.9
## Dec 2020      246063.4
```

## Forecasts from HoltWinters




```
## Growth
```

```
year_2019 <- window(ts, 2019,c(2019,12))
sum_year_2019 = sum(c(year_2019))
year_2020 = (as.data.frame(ts_forcaste2))[1][c(1:11),]
sum_year_2020_HW <-sum(c(year_2020),window(ts, 2020))
growth_HW <- growth(sum_year_2020_HW,sum_year_2019)
growth_HW
```

```
## [1] 0.06859268
```

We can investigate whether the predictive model can be improved upon by checking whether the in-sample forecast errors show non-zero autocorrelations at lags 1-20, by making a correlogram and carrying out the Ljung-Box test:




```
##
## Box-Ljung test
##
## data: ts_forcaste2$residuals
## X-squared = 12.859, df = 20, p-value = 0.8834
```

The correlogram shows that the autocorrelations for the in-sample forecast errors do not exceed the significance bounds for lags 1-20. Furthermore, the p-value for Ljung-Box test is 0.9, indicating that there is no evidence of non-zero autocorrelations at lags 1-20.

We can check whether the forecast errors have constant variance over time, and are normally distributed with mean zero, by making a time plot of the forecast errors and a histogram (with overlaid normal curve):

```
plot.ts(ts_forcaste2$residuals)
```






```
# plotForecastErrors(ts_forcaste2$residuals)
```


From the time plot, it appears plausible that the forecast errors have constant variance over time. From the histogram of forecast errors, it seems plausible that the forecast errors are normally distributed with mean zero.

Thus, there is little evidence of autocorrelation at lags 1-20 for the forecast errors, and the forecast errors appear to be normally distributed with mean zero and constant variance over time. This suggests that Holt-Winters exponential smoothing provides an adequate predictive model of the log of total productivity, which probably cannot be improved upon. Furthermore, the assumptions upon which the prediction intervals were based are probably valid.

```
plot.ts(ts)
```




```
ts_diff1 <- diff(ts, differences = 1)
plot.ts(ts_diff1)
```



The time series of differences (above) does appear to be stationary in mean and variance, as the level of the series stays roughly constant over time, and the variance of the series appears roughly constant over time

```
acf(ts_diff1, lag.max=20) # plot a correlogram
```

V1




We see from the correlogram that the autocorrelation exceeds the significance bound 3 times but all the others do not exceed

```
acf(ts_diff1, lag.max=20, plot=FALSE) # get the autocorrelation values
```

```
##
## Autocorrelations of series 'ts_diff1', by lag
##
## 0.0000 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333
## 1.000 -0.390 -0.066 0.041 0.049 -0.091 0.068 -0.188 0.113 0.079 -0.159
## 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
## -0.199 0.505 -0.141 -0.154 0.109 0.003 -0.035 0.029 -0.129 0.046
```

```
pacf(ts_diff1, lag.max=20) # plot a partial correlogram
```

### Series ts\_diff1



```
pacf(ts_diff1, lag.max=20, plot=FALSE) # get the partial autocorrelation values
```

```
##
## Partial autocorrelations of series 'ts_diff1', by lag
##
## 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333 0.9167
## -0.390 -0.258 -0.116  0.005 -0.074  0.015 -0.222 -0.078  0.060 -0.110 -0.382
## 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
##  0.270  0.216 -0.110  0.006  0.046 -0.042 -0.082  0.025 -0.054
```

**Arima, 0,1,0**

```
ts_arma = Arima(ts, order=c(1,1,1),seasonal = list(order = c(1,1,1)))
ts_arma
```

```
## Series: ts
## ARIMA(1,1,1)(1,1,1)[12]
##
## Coefficients:
##          ar1          ma1          sar1          sma1
##      0.0199   -0.6250   -0.1612   -0.7249
## s.e.  0.1551    0.1227    0.1207    0.1065
```


```
##
## sigma^2 estimated as 178190944: log likelihood=-1314.33
## AIC=2638.66 AICc=2639.19 BIC=2652.6
```

```
ts_arma_forecast = forecast(ts_arma,h = 11)
ts_arma_forecast
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Feb 2020	217750.2	200638.1	234862.4	191579.4	243921.0
## Mar 2020	226899.4	208501.1	245297.6	198761.7	255037.0
## Apr 2020	218520.2	198990.3	238050.1	188651.8	248388.7
## May 2020	214394.1	193795.9	234992.3	182891.8	245896.4
## Jun 2020	235437.7	213824.0	257051.5	202382.3	268493.1
## Jul 2020	233673.6	211089.9	256257.3	199134.8	268212.4
## Aug 2020	236843.1	213329.5	260356.8	200882.1	272804.1
## Sep 2020	237924.8	213516.6	262332.9	200595.7	275253.9
## Oct 2020	246072.2	220801.1	271343.3	207423.4	284721.0
## Nov 2020	244826.2	218720.7	270931.6	204901.3	284751.0
## Dec 2020	246327.3	219413.3	273241.3	205165.9	287488.8

```
forecast::plot.forecast(ts_arma_forecast)
```

### Forecasts from ARIMA(1,1,1)(1,1,1)[12]



## Growth

```
# this_year_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1]

# growth_ARIMA <- growth(sum(c(this_year,as.numeric(this_year_predict_ARIMA$'Point Forecast'))), sum(la
# growth_ARIMA

# year_2019_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1][c(1),]
# sum_year_2019 = sum(c(year_2019))
year_2020 = (as.data.frame(ts_arma_forecast))[1][c(1:11),]
sum_year_2020_ARIMA = sum(c(year_2020>window(ts,2020)))
growth_ARIMA <- growth(sum_year_2020_ARIMA,sum_year_2019)
growth_ARIMA
```


```
## [1] 0.04534847
```

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arma_forecast$residuals, lag.max=20)
```

### Series ts\_arma\_forecast\$residuals




```
Box.test(ts_arima_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: ts_arima_forecast$residuals  
## X-squared = 14.401, df = 20, p-value = 0.8096
```

p value too high to reject


```
plot.ts(ts_arima_forecast$residuals) # make time plot of forecast errors
```



```
plotForecastErrors(ts_arima_forecast$residuals)
```



## Histogram of forecasterrors




## A model chosen automatically

```
fit3 <- auto.arima(ts)
fit3
```

```
## Series: ts
## ARIMA(1,0,1)(0,1,1)[12] with drift
##
## Coefficients:
##          ar1          ma1          sma1          drift
##          0.9348 -0.5736 -0.8165 1208.7731
## s.e.  0.0435  0.0962  0.0938  177.2516
##
## sigma^2 estimated as 172154591:  log likelihood=-1323.51
## AIC=2657.03  AICc=2657.55  BIC=2671.01
```

```
fit_forecast = forecast(fit3,h=11)
plot(fit_forecast)
```

## Forecasts from ARIMA(1,0,1)(0,1,1)[12] with drift



```
# str(fit)
```

## Growth

```
# year_2019 <- window(ts, 2019)
# year_2019_predict_HW <- (as.data.frame(ts_forcaste2))[1][c(1),]
# sum_year_2019 = sum(c(year_2019))
# year_2020 = (as.data.frame(ts_forcaste2))[1][c(2:13),]

# year_2019_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(1),]
# year_2019_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(1),]
# year_2019_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(1),]

# sum_year_2019 = sum(c(year_2019))
# sum_year_2019_low = sum(c(year_2019, year_2019_predict_auto.arima_95_low))
# sum_year_2019_high = sum(c(year_2019, year_2019_predict_auto.arima_95_high))

year_2020_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(1:11),]
# year_2020_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(2:13),]
# year_2020_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(2:13),]
```

```

sum_year_2020_auto.ARIMA = sum(c(year_2020_predict_auto.arima>window(ts,2020)))

growth_auto.arima <- growth(sum_year_2020_auto.ARIMA,sum_year_2019)
# growth_auto.arima_95_low <- growth(sum(year_2020_predict_auto.arima_95_low),sum_year_2019_low)
# growth_auto.arima_95_high <- growth(sum(year_2020_predict_auto.arima_95_high),sum_year_2019_high)

growth_auto.arima

## [1] 0.03690604

# growth_auto.arima_95_low
# growth_auto.arima_95_high

```