

tsf_export

Kevork Sulahian

October 28, 2019

```
library(readxl)
library(forecast)
```

```
# library(readxl)
```

```
df <- read_xlsx("Export_for_TS.xlsx")
```

```
## New names:
## * `` -> ...2
```

```
df = df[1,]
df = df[-c(1,2)]
```

```
df2 = read_xlsx('export_19xlsx.xlsx')
```


```
## New names:
## * `` -> ...1
```

```
df = t(df)
df2$...1 = c(paste0(2019,"-",1:9))
rownames(df2) =df2$...1
```

```
## Warning: Setting row names on a tibble is deprecated.
```


```
df2$...1 = NULL
colnames(df2) = "V1"
df = as.data.frame(df)
```

```
df3 = rbind(df,df2)
ts = ts(df3,start=c(2009,1), frequency = c(12))
```



In this case, it appears that an additive model is not appropriate for describing this time series, since the size of the seasonal fluctuations and random fluctuations seem to increase with the level of the time series. Thus, we may need to transform the time series in order to get a transformed time series that can be described using an additive model. For example, we can transform the time series by calculating the natural log of the original data:

```
log_ts <- log(ts)
plot.ts(log_ts)
```



Decomposing Time Series

Decomposing a time series means separating it into its constituent components, which are usually a trend component and an irregular component, and if it is a seasonal time series, a seasonal component.

Decomposing Seasonal Data

A seasonal time series consists of a trend component, a seasonal component and an irregular component. Decomposing the time series means separating the time series into these three components: that is, estimating these three components.

```
ts_components <- decompose(ts)
```


we can print out the estimated values of the seasonal component

```
ts_components$seasonal
```

| ## | Jan | Feb | Mar | Apr | May |
|---------|-------------|------------|----------|------------|------------|
| ## 2009 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## 2010 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## 2011 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## 2012 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## 2013 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |

| | | | | | | |
|----|------|-------------|------------|-----------|------------|------------|
| ## | 2014 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## | 2015 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## | 2016 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## | 2017 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## | 2018 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## | 2019 | -35837.3733 | -5078.4271 | 614.7571 | -5025.8461 | -7561.9849 |
| ## | | Jun | Jul | Aug | Sep | Oct |
| ## | 2009 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2010 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2011 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2012 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2013 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2014 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2015 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2016 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2017 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2018 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | 8951.2410 |
| ## | 2019 | 8910.0983 | 3476.2085 | 2185.4111 | 8757.6873 | |
| ## | | Nov | Dec | | | |
| ## | 2009 | 8780.2368 | 11827.9913 | | | |
| ## | 2010 | 8780.2368 | 11827.9913 | | | |
| ## | 2011 | 8780.2368 | 11827.9913 | | | |
| ## | 2012 | 8780.2368 | 11827.9913 | | | |
| ## | 2013 | 8780.2368 | 11827.9913 | | | |
| ## | 2014 | 8780.2368 | 11827.9913 | | | |
| ## | 2015 | 8780.2368 | 11827.9913 | | | |
| ## | 2016 | 8780.2368 | 11827.9913 | | | |
| ## | 2017 | 8780.2368 | 11827.9913 | | | |
| ## | 2018 | 8780.2368 | 11827.9913 | | | |
| ## | 2019 | | | | | |


Decomposition of additive time series



The plot above shows the original time series (top), the estimated trend component (second from top), the estimated seasonal component (third from top), and the estimated irregular component (bottom)

Seasonally Adjusting

```
ts_seasonall <- ts - ts_components$seasonal
```



Holt-Winters Exponential Smoothing

```
ts_forcaste <- HoltWinters(ts)
ts_forcaste
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = ts)
##
## Smoothing parameters:
##  alpha: 0.3721975
##  beta : 0.008388046
##  gamma: 0.3879005
##
## Coefficients:
##           [,1]
## a  234236.22519
## b   1800.90011
## s1  10884.32142
## s2   2422.24425
## s3   6639.03969
## s4  -45793.46458
```


```
## s5    -7778.48158
## s6     -48.50151
## s7   -8403.94186
## s8  -12451.28939
## s9    4905.92488
## s10   6841.04904
## s11   6190.87742
## s12   4206.41639
```

```
#
```

The value of alpha (0.35) is relatively low, indicating that the estimate of the level at the current time point is based upon both recent observations and some observations in the more distant past. The value of beta is 0.01, indicating that the estimate of the slope b of the trend component is updated but doesn't have much effect over the time series, and instead is set equal to its initial value. This makes good intuitive sense, as the level changes quite a bit over the time series, but the slope b of the trend component remains roughly the same. In contrast, the value of gamma (0.38) is high, indicating that the estimate of the seasonal component at the current time point is not just based upon very recent observations

```
ts_forcaste$SSE
```


```
## [1] 21744166076
```



```
ts_forcaste2 = forecast::forecast.HoltWinters(ts_forcaste, h= 15)
(as.data.frame(ts_forcaste2))[1]
```

```
##          Point Forecast
## Oct 2019      246921.4
## Nov 2019      240260.3
## Dec 2019      246278.0
## Jan 2020      195646.4
## Feb 2020      235462.2
## Mar 2020      244993.1
## Apr 2020      238438.6
## May 2020      236192.1
## Jun 2020      255350.3
## Jul 2020      259086.3
## Aug 2020      260237.0
## Sep 2020      260053.4
## Oct 2020      268532.2
## Nov 2020      261871.1
## Dec 2020      267888.8
```

Forecasts from HoltWinters



```
## Growth
```


```
year_2019 <- window(ts, 2019)
year_2019_predict_HW <- (as.data.frame(ts_forcaste2))[1][c(1:3),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_HW))
year_2020 = (as.data.frame(ts_forcaste2))[1][c(4:15),]
```



```
growth_HW <- growth(sum(year_2020),sum_year_2019)
growth_HW
```

```
## [1] 0.1369446
```

We can investigate whether the predictive model can be improved upon by checking whether the in-sample forecast errors show non-zero autocorrelations at lags 1-20, by making a correlogram and carrying out the Ljung-Box test:




```
##
## Box-Ljung test
##
## data: ts_forcaste2$residuals
## X-squared = 10.37, df = 20, p-value = 0.961
```

The correlogram shows that the autocorrelations for the in-sample forecast errors do not exceed the significance bounds for lags 1-20. Furthermore, the p-value for Ljung-Box test is 0.9, indicating that there is no evidence of non-zero autocorrelations at lags 1-20.


We can check whether the forecast errors have constant variance over time, and are normally distributed with mean zero, by making a time plot of the forecast errors and a histogram (with overlaid normal curve):

```
plot.ts(ts_forcaste2$residuals)
```



```
plotForecastErrors(ts_forcaste2$residuals)
```


Histogram of forecasterrors




From the time plot, it appears plausible that the forecast errors have constant variance over time. From the histogram of forecast errors, it seems plausible that the forecast errors are normally distributed with mean zero.

Thus, there is little evidence of autocorrelation at lags 1-20 for the forecast errors, and the forecast errors appear to be normally distributed with mean zero and constant variance over time. This suggests that Holt-Winters exponential smoothing provides an adequate predictive model of the log of total productivity, which probably cannot be improved upon. Furthermore, the assumptions upon which the prediction intervals were based are probably valid.

```
plot.ts(ts)
```




```
ts_diff1 <- diff(ts, differences = 1)
plot.ts(ts_diff1)
```



The time series of differences (above) does appear to be stationary in mean and variance, as the level of the series stays roughly constant over time, and the variance of the series appears roughly constant over time

```
acf(ts_diff1, lag.max=20) # plot a correlogram
```

V1




We see from the correlogram that the autocorrelation exceeds the significance bound 3 times but all the others do not exceed

```
acf(ts_diff1, lag.max=20, plot=FALSE) # get the autocorrelation values
```

```
##
## Autocorrelations of series 'ts_diff1', by lag
##
## 0.0000 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500
## 1.000 -0.442 -0.056 0.038 0.049 -0.085 0.098 -0.200 0.126 0.064
## 0.8333 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833
## -0.108 -0.213 0.508 -0.190 -0.157 0.143 -0.010 -0.043 0.027 -0.111
## 1.6667
## 0.039
```

```
pacf(ts_diff1, lag.max=20) # plot a partial correlogram
```



```
pacf(ts_diff1, lag.max=20, plot=FALSE) # get the partial autocorrelation values
```

```
##
## Partial autocorrelations of series 'ts_diff1', by lag
##
## 0.0833 0.1667 0.2500 0.3333 0.4167 0.5000 0.5833 0.6667 0.7500 0.8333
## -0.442 -0.312 -0.182 -0.044 -0.097 0.038 -0.211 -0.095 0.036 -0.053
## 0.9167 1.0000 1.0833 1.1667 1.2500 1.3333 1.4167 1.5000 1.5833 1.6667
## -0.379 0.255 0.225 -0.117 0.047 0.089 -0.019 -0.126 0.028 -0.068
```

Arima, 0,1,0

```
ts_arima = Arima(ts, order=c(0,1,0),seasonal = list(order = c(0,1,0)))
ts_arima
```


```
## Series: ts
## ARIMA(0,1,0)(0,1,0)[12]
##
## sigma^2 estimated as 397639845: log likelihood=-1313.06
## AIC=2628.12 AICc=2628.15 BIC=2630.87
```

```
ts_arima_forecast = forecast(ts_arima,h = 15)
ts_arima_forecast
```

| ## | Point Forecast | Lo 80 | Hi 80 | Lo 95 | Hi 95 |
|-------------|----------------|----------|----------|----------|----------|
| ## Oct 2019 | 291503.5 | 265948.2 | 317058.8 | 252420.1 | 330587.0 |
| ## Nov 2019 | 263011.9 | 226871.2 | 299152.5 | 207739.5 | 318284.2 |
| ## Dec 2019 | 268333.3 | 224070.2 | 312596.4 | 200638.8 | 336027.9 |
| ## Jan 2020 | 205722.4 | 154611.8 | 256833.0 | 127555.5 | 283889.3 |
| ## Feb 2020 | 252922.4 | 195779.0 | 310065.8 | 165529.1 | 340315.7 |
| ## Mar 2020 | 266822.4 | 204225.0 | 329419.9 | 171087.9 | 362556.9 |
| ## Apr 2020 | 268222.4 | 200609.4 | 335835.4 | 164817.3 | 371627.5 |
| ## May 2020 | 268622.4 | 196341.1 | 340903.7 | 158077.7 | 379167.1 |
| ## Jun 2020 | 272622.4 | 195956.5 | 349288.3 | 155372.0 | 389872.8 |
| ## Jul 2020 | 295822.4 | 215009.4 | 376635.4 | 172229.6 | 419415.2 |
| ## Aug 2020 | 299922.4 | 215165.1 | 384679.8 | 170297.2 | 429547.6 |
| ## Sep 2020 | 303122.4 | 214596.2 | 391648.6 | 167733.3 | 438511.5 |
| ## Oct 2020 | 351825.9 | 249604.7 | 454047.1 | 195492.1 | 508159.8 |
| ## Nov 2020 | 323334.3 | 209047.5 | 437621.1 | 148547.7 | 498120.8 |
| ## Dec 2020 | 328655.7 | 203460.8 | 453850.6 | 137186.6 | 520124.8 |

```
forecast::plot.forecast(ts_arima_forecast)
```

Forecasts from ARIMA(0,1,0)(0,1,0)[12]



Growth

```
this_year_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1]

# growth_ARIMA <- growth(sum(c(this_year,as.numeric(this_year_predict_ARIMA$`Point Forecast`))), sum(la
# growth_ARIMA


year_2019_predict_ARIMA <- (as.data.frame(ts_arma_forecast))[1][c(1:3),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_ARIMA))
year_2020 = (as.data.frame(ts_arma_forecast))[1][c(4:15),]
growth_ARIMA <- growth(sum_year_2019, sum(year_2020))
-growth_ARIMA
```

```
## [1] 0.2105728
```

As in the case of exponential smoothing models, it is a good idea to investigate whether the forecast errors of an ARIMA model are normally distributed with mean zero and constant variance, and whether there are correlations between successive forecast errors.

For example, we can make a correlogram of the forecast errors for our ARIMA(0,1,1) model, and perform the Ljung-Box test for lags 1-20, by typing:

```
acf(ts_arma_forecast$residuals, lag.max=20)
```




```
Box.test(ts_arma_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: ts_arma_forecast$residuals  
## X-squared = 126.86, df = 20, p-value < 2.2e-16
```


p value too high to reject

```
plot.ts(ts_arma_forecast$residuals) # make time plot of forecast errors
```



```
plotForecastErrors(ts_arma_forecast$residuals)
```

Histogram of forecasterrors




A model chosen automatically

```
fit3 <- auto.arima(ts)
fit3
```

```
## Series: ts
## ARIMA(1,0,1)(0,1,1)[12] with drift
##
## Coefficients:
##          ar1          ma1          sma1          drift
##          0.9482   -0.5785   -0.8459   1269.853
## s.e.    0.0409    0.0914    0.1186    203.636
##
## sigma^2 estimated as 162448032:  log likelihood=-1277.32
## AIC=2564.64   AICc=2565.18   BIC=2578.45
```

```
fit_forecast = forecast(fit3,h=15)
plot(fit_forecast)
```

Forecasts from ARIMA(1,0,1)(0,1,1)[12] with drift



```
# str(fit)
```

Growth

```
year_2019 <- window(ts, 2019)
year_2019_predict_HW <- (as.data.frame(ts_forcaste2))[1][c(1:3),]
sum_year_2019 = sum(c(year_2019,year_2019_predict_HW))
year_2020 = (as.data.frame(ts_forcaste2))[1][c(4:15),]
growth_HW <- growth(sum(year_2020),sum_year_2019)
growth_HW
```

```
## [1] 0.1369446
```

```
year_2019_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(1:3),]
year_2019_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(1:3),]
year_2019_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(1:3),]

sum_year_2019 = sum(c(year_2019,year_2019_predict_auto.arima))
sum_year_2019_low = sum(c(year_2019,year_2019_predict_auto.arima_95_low))
sum_year_2019_high = sum(c(year_2019,year_2019_predict_auto.arima_95_high))

year_2020_predict_auto.arima <- (as.data.frame(fit_forecast))[1][c(4:15),]
```

```

year_2020_predict_auto.arima_95_low <- (as.data.frame(fit_forecast))[4][c(4:15),]
year_2020_predict_auto.arima_95_high <- (as.data.frame(fit_forecast))[5][c(4:15),]

growth_auto.arima <- growth(sum(year_2020_predict_auto.arima),sum_year_2019)
growth_auto.arima_95_low <- growth(sum(year_2020_predict_auto.arima_95_low),sum_year_2019_low)
growth_auto.arima_95_high <- growth(sum(year_2020_predict_auto.arima_95_high),sum_year_2019_high)

growth_auto.arima

```

```
## [1] 0.08897661
```

```
growth_auto.arima_95_low
```

```
## [1] -0.03516189
```

```
growth_auto.arima_95_high
```

```
## [1] 0.2057449
```